

# ACM SIG AI - SEM 4 Recruitment Tasks

---

## Task 1: Neural Network for Image Classification with Preprocessing

### Objective:

Develop a deep learning pipeline that preprocesses images and classifies them using a neural network.

### Requirements:

**Preprocessing, Model Development, Training and Evaluation**

### Bonus:

- Implement Grad-CAM to visualize which parts of an image influenced the classification.
  - Compare CNN performance with a Vision Transformer (ViT) model.
- 

## Task 2: Named Entity Recognition (NER) Using Transformers

### Objective:

Fine-tune a transformer model for Named Entity Recognition (NER) to identify entities like names, locations, and organizations.

### Dataset:

- Use the **CoNLL-2003** dataset or a custom dataset with labeled entities.

### Requirements:

### Data Preprocessing:

1. Tokenize the text using a transformer-based tokenizer (e.g., BERT tokenizer).
2. Convert labels to a format compatible with the model (e.g., IOB format).

3. Apply padding and truncation to handle different sequence lengths.

## Model Development and Evaluation

### Bonus:

- Build a **Streamlit app** where users can enter text and get highlighted named entities in real time.
- 

## Task 3: Chatbot Using Retrieval-Augmented Generation (RAG)

### Objective:

Develop a **Retrieval-Augmented Generation (RAG) chatbot** that retrieves relevant documents before generating a response using a large language model (LLM).

### Dataset:

- Use **Wikipedia dumps, product manuals, FAQs**, or any domain-specific knowledge base.

### Requirements:

#### Data Processing & Indexing:

1. Store documents in an efficient retrieval system using **FAISS** or **ChromaDB**.
2. Preprocess documents: remove stopwords, lowercase, tokenize.
3. Convert text into embeddings using **Sentence Transformers** or OpenAI's embeddings API.

#### Model Development:

1. Implement **retrieval** using FAISS or ChromaDB to fetch relevant passages.
2. Use a **pre-trained LLM** (GPT-3.5, LLaMA, or T5) to generate responses based on retrieved knowledge.
3. Fuse the retrieved information with the generated response using a prompt-engineering strategy.

#### Evaluation:

1. Use **BERTScore** or **manual evaluation** to assess response quality.
2. Log chatbot interactions and analyze failure cases.

3. Compare performance with and without retrieval (i.e., knowledge-augmented vs. vanilla LLM).

**Bonus:**

- Add **speech-to-text (STT) and text-to-speech (TTS)** for a voice-based chatbot experience.
  - Deploy the chatbot using **Gradio or Streamlit**.
- 

## Task 4: AI-Powered Code Assistant for Code Editors

**Objective:**

Build an AI-powered coding assistant that integrates with a **code editor (VS Code, Jupyter, etc.)**, offering real-time code suggestions and simplifications.

**Requirements:**

**Core Features:**

1. Implement an **autocompletion engine** using OpenAI Codex, CodeT5, or StarCoder.
2. Provide **code refactoring suggestions** for improving efficiency.
3. Offer **error detection and auto-fix suggestions**.
4. Summarize code logic in natural language.

**Deployment:**

1. Package as a **VS Code extension** or a **Jupyter Notebook plugin**.
2. Use **Streamlit or Flask** to build a web-based interactive demo.

**Evaluation:**

1. Compare AI suggestions with **ground truth implementations**.
2. Measure **latency (response time)** and **user satisfaction (manual review, Likert scale)**.

**Bonus:**

- Implement **context-aware suggestions** based on previous lines of code.
  - Add **multi-language support (Python, Java, JavaScript, C++)**.
-

# Bonus Task: End-to-End Facial Recognition Attendance System

## Objective:

Develop a **facial recognition-based attendance system** that detects and logs attendance in real time.

## Dataset:

- Use **LFW (Labeled Faces in the Wild)** or collect a **custom dataset**.

## Requirements:

### Preprocessing:

1. Detect and align faces using **MTCNN** or **OpenCV**.
2. Extract facial embeddings using **FaceNet** or **DeepFace**.

### Model Development:

1. Train a **Siamese network** or use **pre-trained FaceNet** for recognition.
2. Store facial embeddings in a **database (SQLite, MongoDB, or PostgreSQL)**.
3. Match incoming faces with stored embeddings to mark attendance.

### Deployment:

1. Build a **Streamlit/Flask app** with a webcam-based interface.
2. Display real-time logs of attendance.

### Evaluation:

1. Use **accuracy, false acceptance rate (FAR), and false rejection rate (FRR)**.
2. Compare recognition time and accuracy against traditional ID-based attendance.

---

## Submission Guidelines

### Code:

- Use **Python** for all implementations.
- Organize tasks in **separate folders** (**Task1**, **Task2**, etc.).
- Use **Jupyter Notebooks** or **Python scripts**.

## Documentation:

- Each task must include a **README.md** explaining:
  - The problem statement
  - Steps taken
  - How to run the code
  - Results summary

## Visualizations:

- Include **graphs, confusion matrices, and evaluation reports** where applicable.

## GitHub Repository:

- Name the repo: **yourname\_ACM\_SIGAI\_Recr\_S4**.
- Make it **public** and ensure proper **folder structuring**.

## Demonstration is Mandatory

**Good luck!**

LOKESH YARRAMALLU, SIG AI LEAD