



**Amrita Vishwa Vidyapeetham**  
Amritapuri Campus

# Introduction to Transformers



# **Transformers : Self Attention**

---

# Attention Is All You Need

---

Ashish Vaswani\*  
Google Brain  
avaswani@google.com

Noam Shazeer\*  
Google Brain  
noam@google.com

Niki Parmar\*  
Google Research  
nikip@google.com

Jakob Uszkoreit\*  
Google Research  
usz@google.com

Llion Jones\*  
Google Research  
llion@google.com

Aidan N. Gomez\* †  
University of Toronto  
aidan@cs.toronto.edu

Lukasz Kaiser\*  
Google Brain  
lukaszkaiser@google.com

Illia Polosukhin\* ‡  
illia.pолосухин@gmail.com

- Google Brain , University of Toronto 2017 [Vaswani et al. 2017].

## Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

# Why Transformers?

- Recurrent architectures rely on sequential processing of input at the encoding step that results in computational inefficiency, as the processing cannot be parallelized
- Transformer architecture completely eliminates sequential processing and recurrent connections. It relies only on **self attention** mechanism to capture global dependencies between input and output
- Significant **parallel processing, shorter training time and higher accuracy** for Machine Translation without any recurrent component

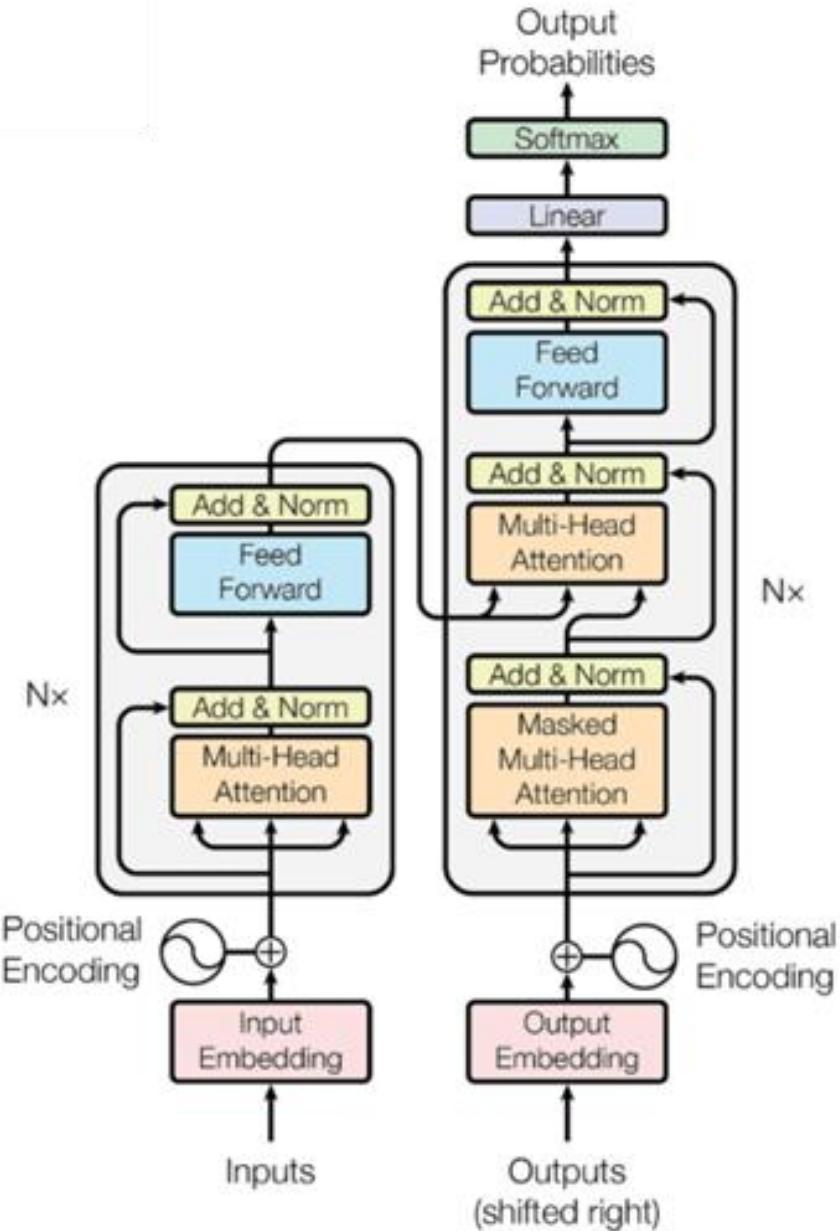
## Challenges of RNN/LSTM

- **Sequential processing** Bottle Neck- training slow for large datasets or long sequence.
- **Issue with Long Range Dependencies-** Understanding limited to immediate context
- **Gradient vanishing and Exploding**
- Large # of **training steps**
- **Resource inefficiency-** Sequential nature, limits scalability and training on massive datasets
- Recurrence **prevents parallel** computation

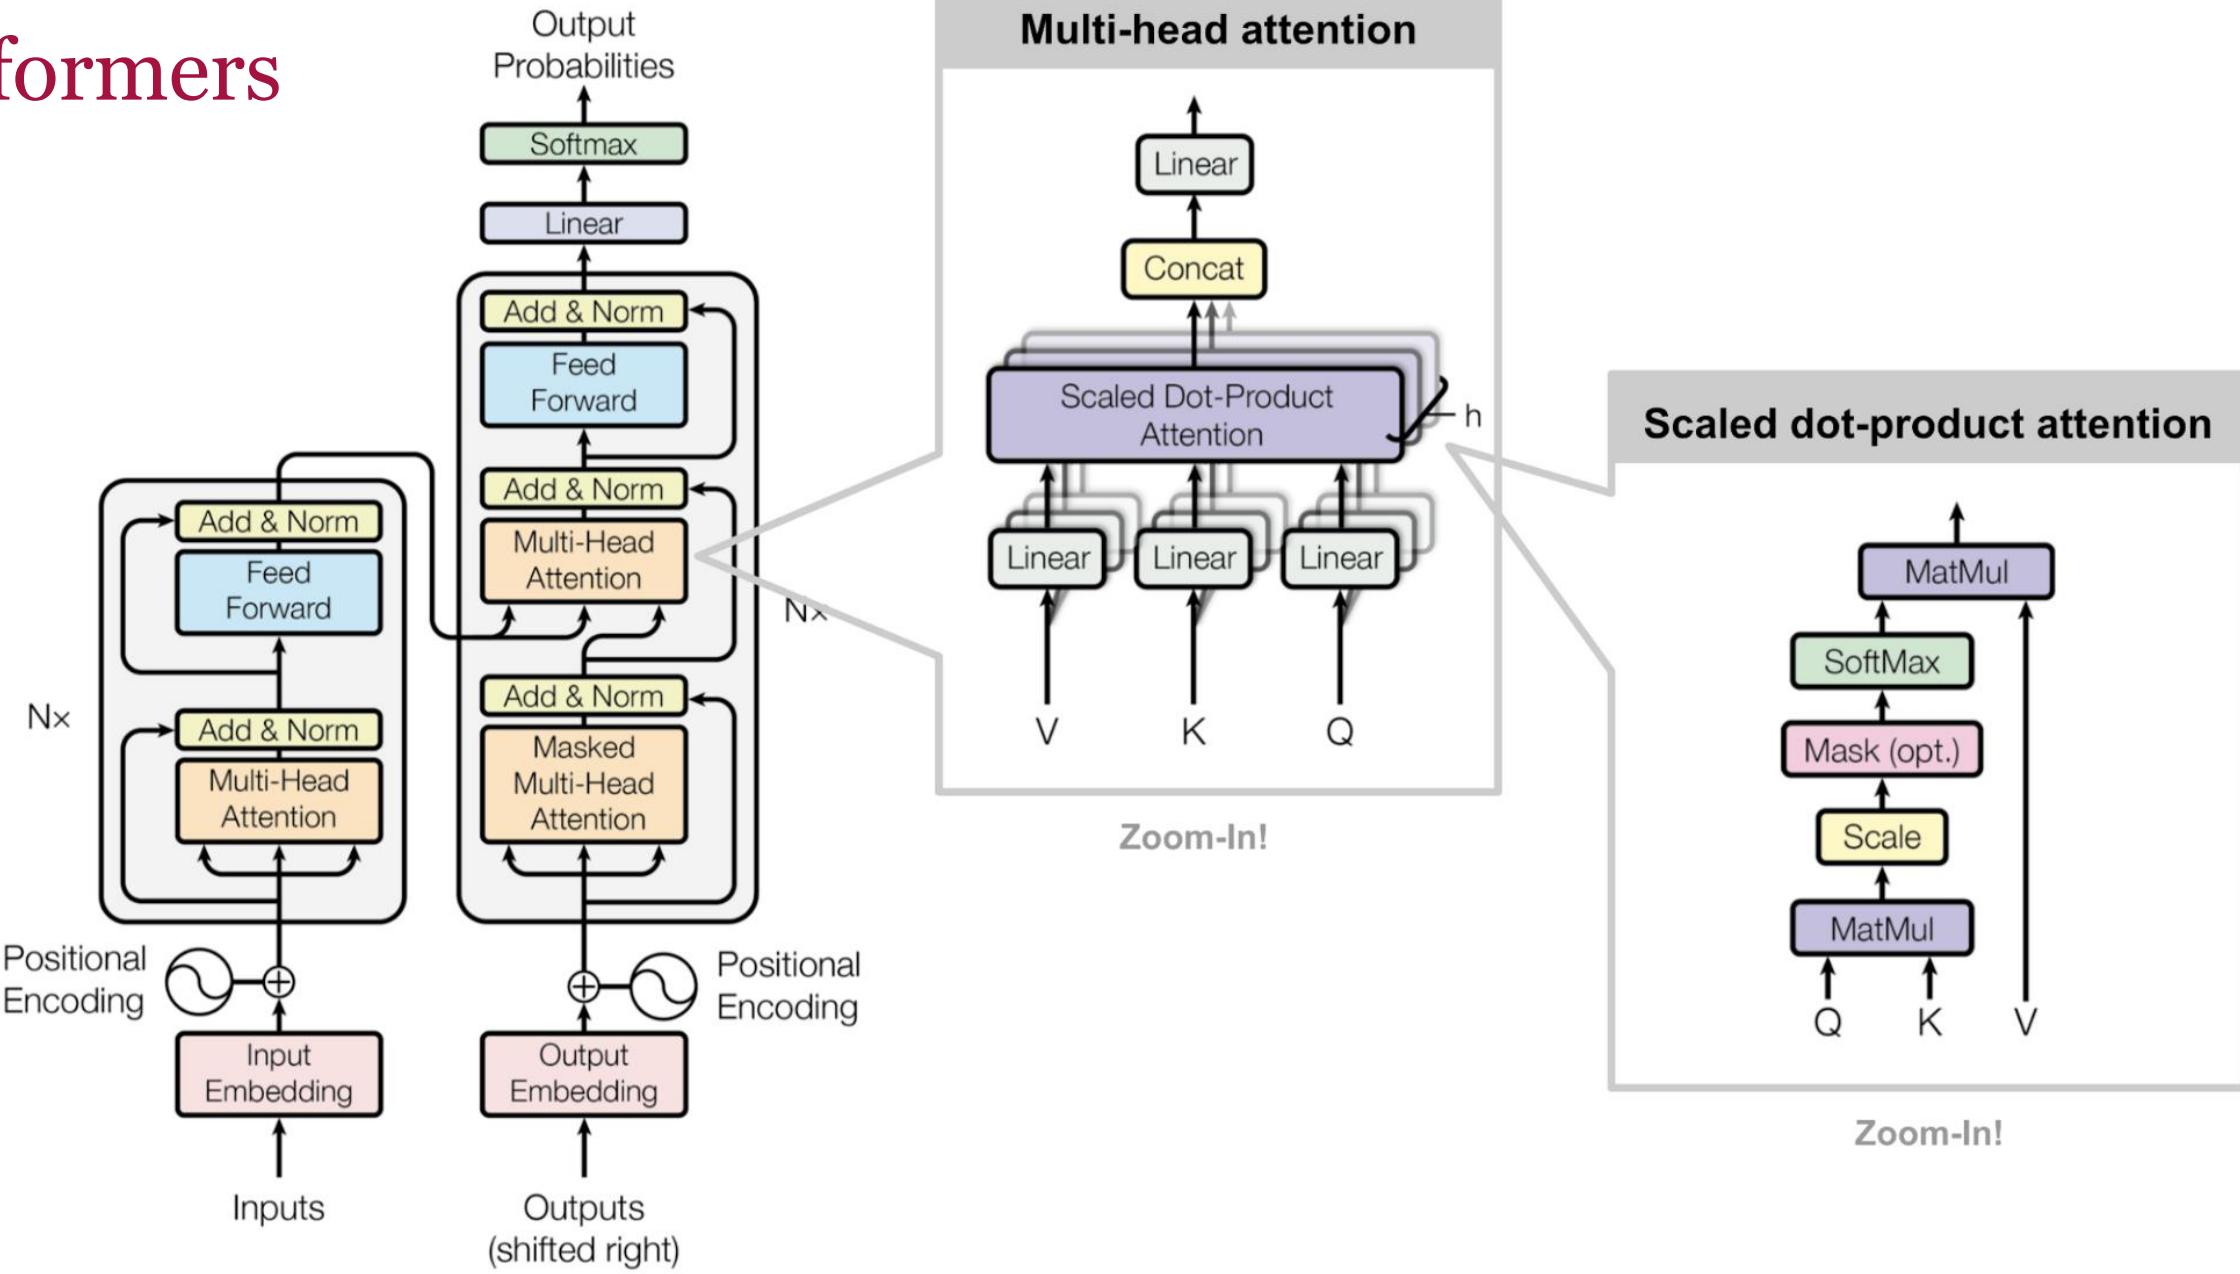
## Transformer Networks

- **Parallelisation** through self attention
- **Faster training**
- **Facilitate long range dependencies-** models dependencies between any 2 tokens regardless of position- retain context across long docs
- No gradient vanishing and explosion
- Scalability and Efficiency
- Fewer training steps
- Flexible Context modeling
- No recurrence and that facilitate parallel computation

# Transformers



# Transformers



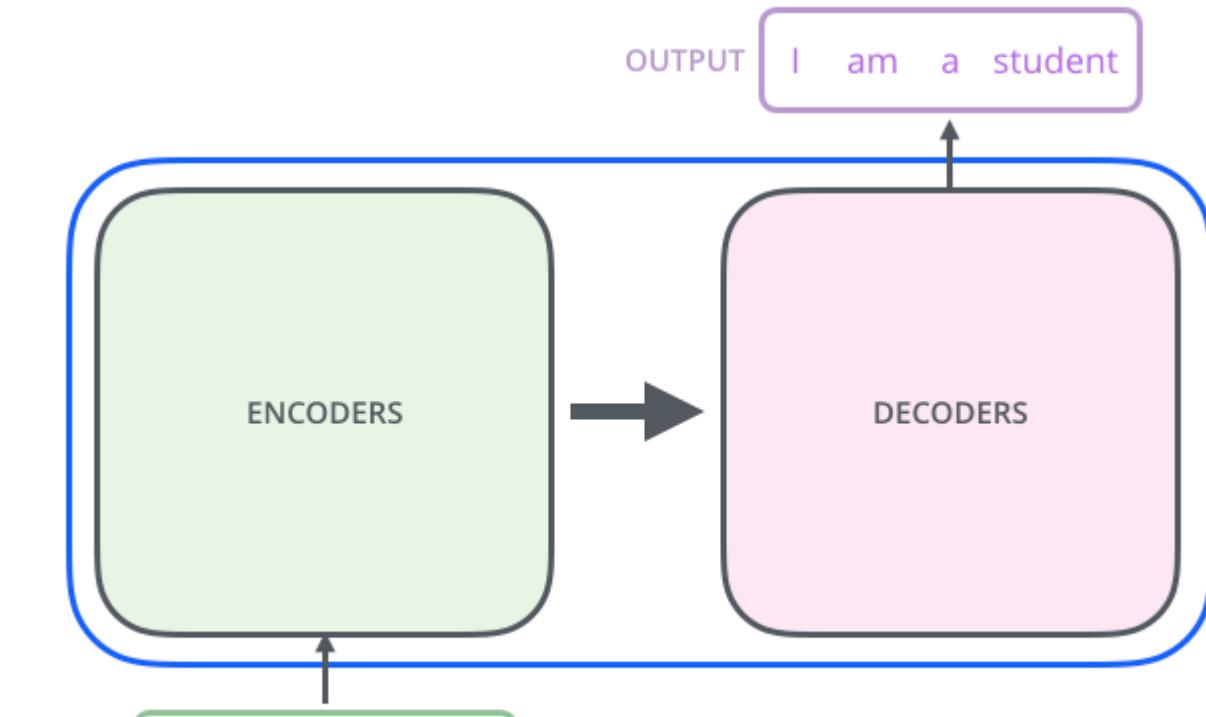
# Transformers – Machine Translation example



In a machine translation application, it would take a sentence in one language, and output its translation in another.

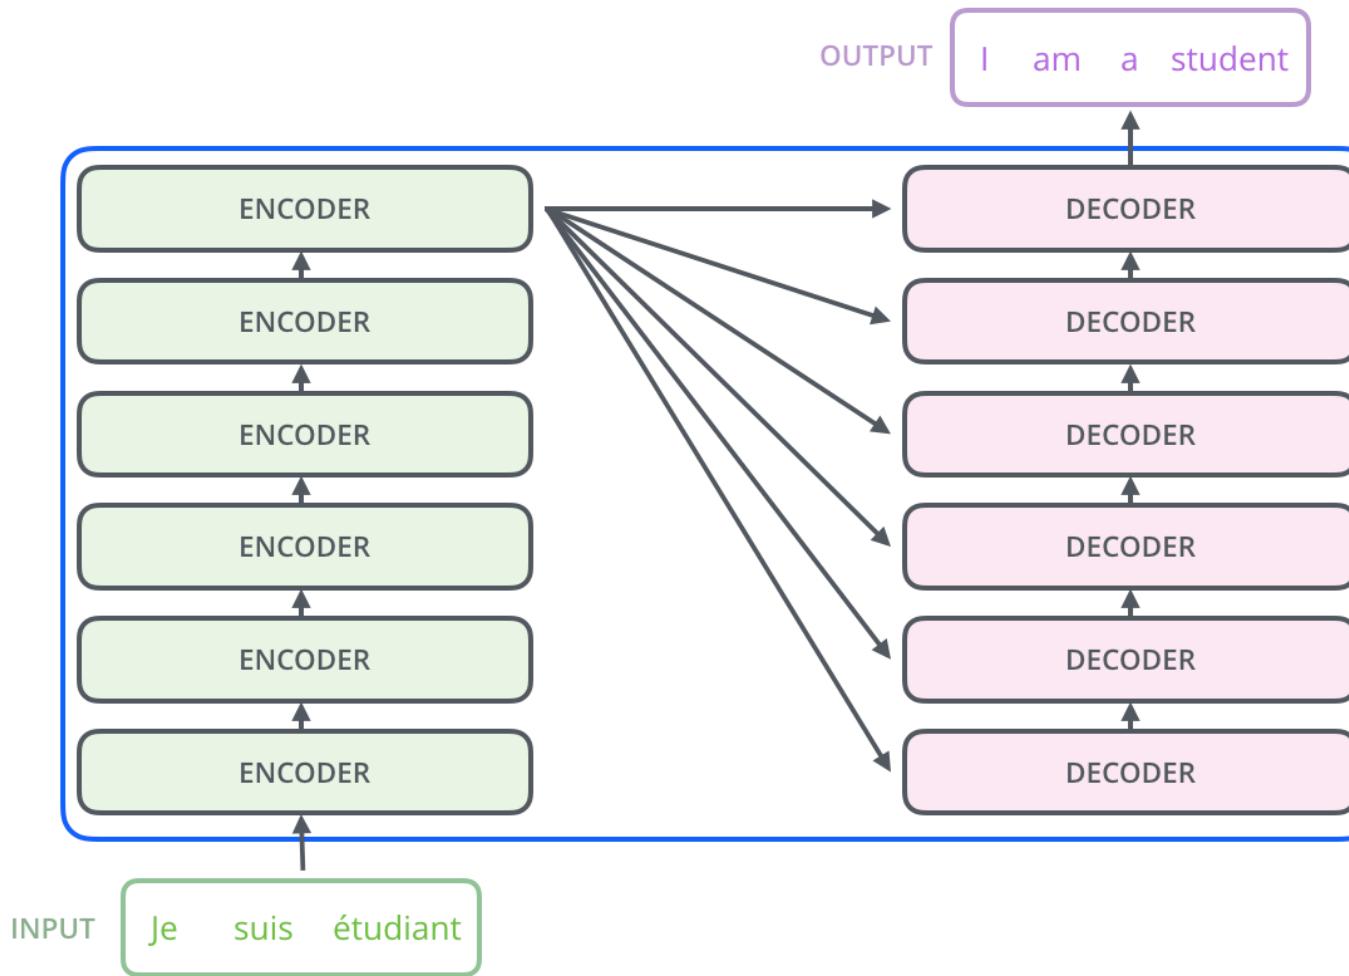
Courtesy : [jalammar.github.io](https://jalammar.github.io)

# Transformers – Machine Translation example



An encoding component, a decoding component, and connections between them.

# Transformers – Machine Translation example

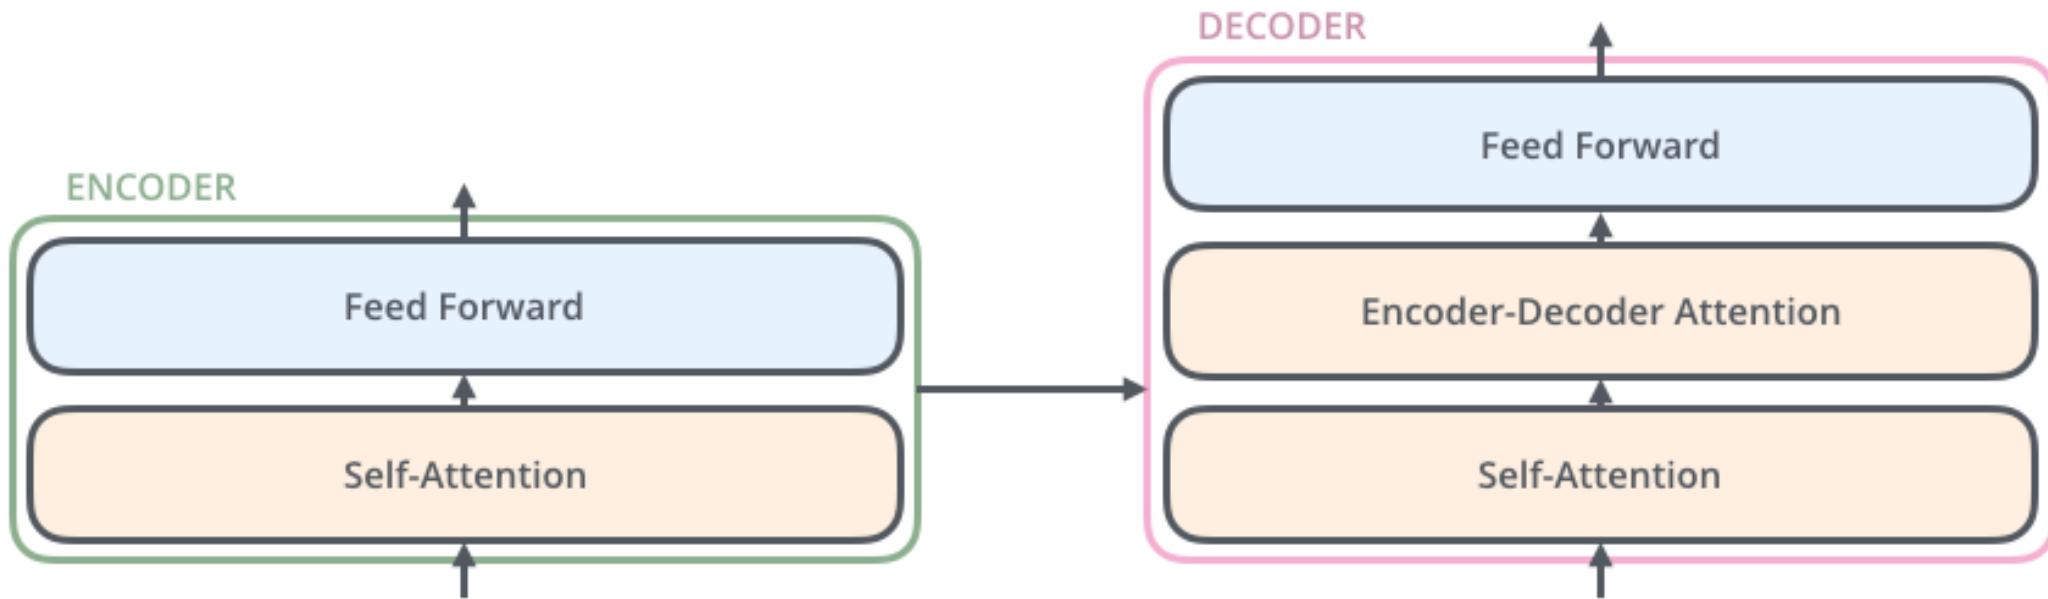


The encoders are all identical in structure (yet they do not share weights).

Each one is broken down into two sub-layers:

The encoding component is a stack of encoders .The decoding component is a stack of decoders of the same number.

# Transformers – Machine Translation example



The encoder's inputs first flow through a self-attention layer – a layer that helps the encoder look at other words in the input sentence as it encodes a specific word.

The outputs of the self-attention layer are fed to a feed-forward neural network.

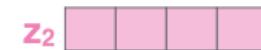
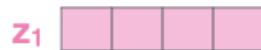
ENCODER #2

ENCODER #1

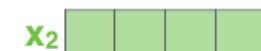
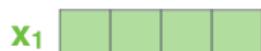


Feed Forward  
Neural Network

Feed Forward  
Neural Network



Self-Attention



Thinking

Machines

# Self-Attention

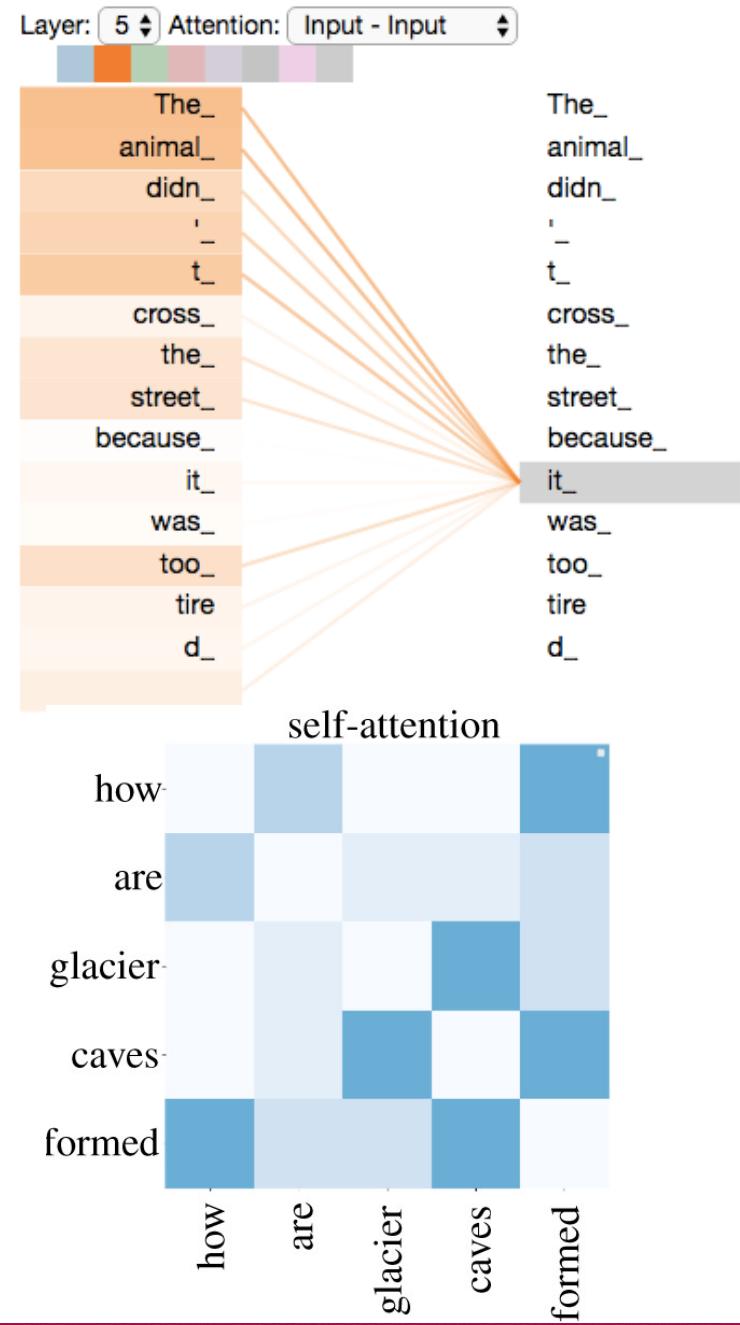
The animal didn't cross the street because it was too tired”

What does “it” in this sentence refer to? Is it referring to the **street** or to the **animal**? It’s a simple question to a human, but not as simple to an algorithm.

When the model is processing the word “it”, self-attention allows it to associate “it” with “animal”.

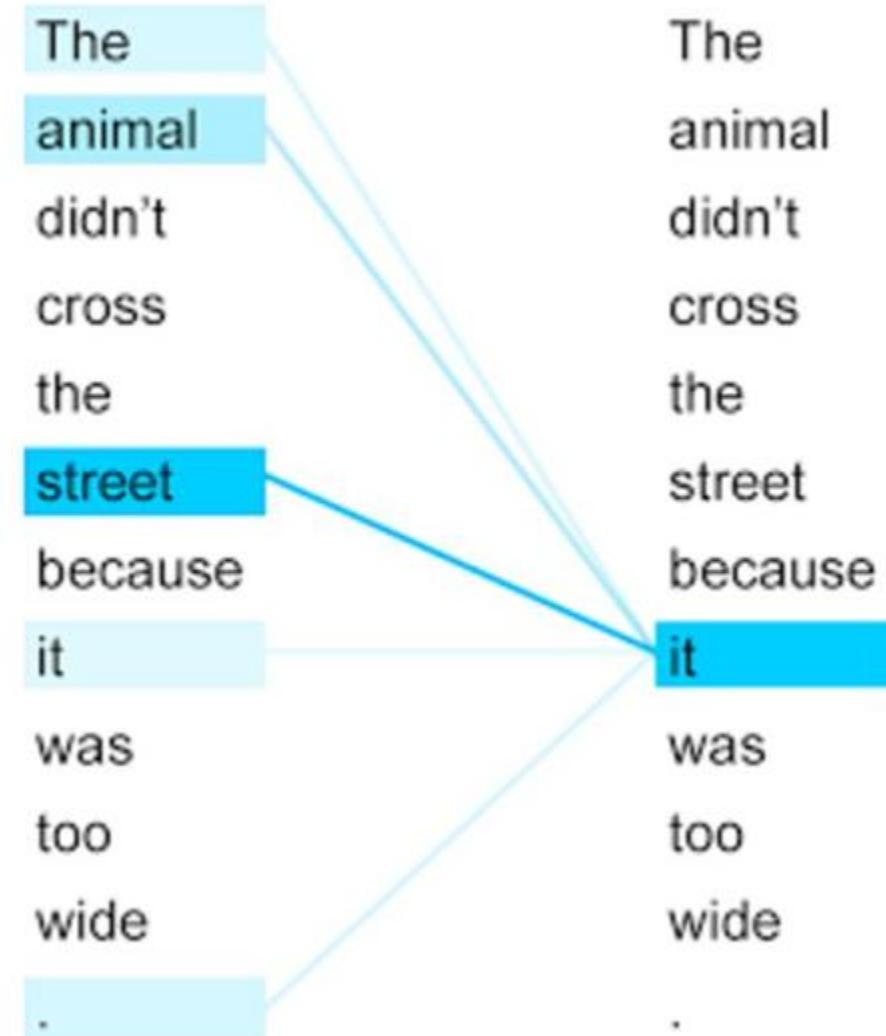
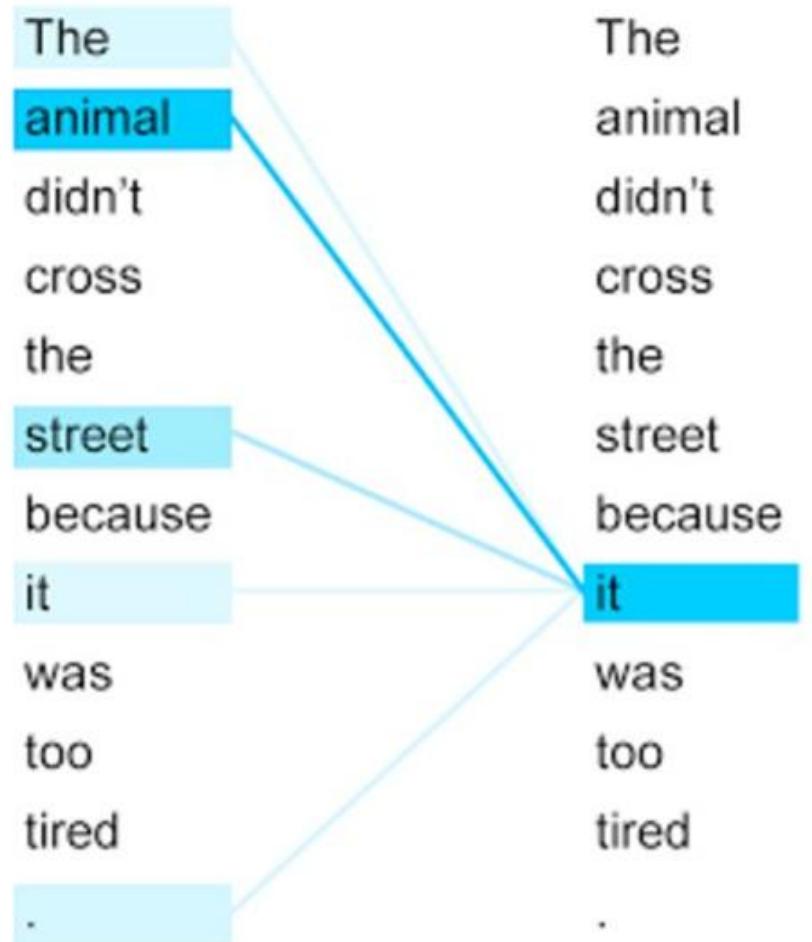
Self-attention is the method the Transformer uses to bake the “understanding” of other relevant words into the one we’re currently processing.

**Self-attention** is a mechanism in transformers that allows each token in a sequence to weigh and attend to all other tokens, capturing contextual relationships dynamically.



# Self attention

**Self-attention** is a mechanism in transformers where each token in a sequence computes weighted dependencies with all other tokens using **query, key, and value projections**, enabling dynamic context-aware representations with parallelized computation.



# Self Attention

- **Query vector (  $q_i$  )**

- what we are looking for in the sequence

- **Key vector (  $k_i$  )**

- what the element is “offering”,

- **Value vector (  $v_i$  )**

This feature vector is the one we want to average over and finally get extracted

- **Score**

To rate which elements we want to pay attention to

Query →

key1

key2

key3

key4

Amrita Vishwa Vidyapeetham Amritapuri

FILTERS

Amrita Vishwa Vidyapeetham, Amritapuri Campus, Kollam, Kerala  
152K views • 2 years ago  
Amrita Vishwa Vidyapeetham  
The Amritapuri campus, nestled in the backdrop of the beautiful village of Vallikavu, provides warmth...

Amrita Vishwa Vidyapeetham, Amritapuri Campus  
17.4K subscribers • 223 videos  
This channel is to publicize the programmes offered at Amrita University, Amritapuri Campus.

Amritavarsham 61  
33K views • 7 years ago  
Amrita Vishwa Vidyapeetham, Amritapuri Campus  
Amma's 61st Birthday Celebration @ Amritapuri.

Unity in Diversity , Amrita University, Amritapuri  
22K views • 6 years ago  
Amrita Vishwa Vidyapeetham, Amritapuri Campus  
Unity in Diversity , Amrita University, Amritapuri.

# Steps in Self Attention

- Query vector (  $q_i$  )
- Key vector (  $k_i$  )
- Value vector (  $v_i$  )

YouTube search results for "Amrita Vishwa Vidyapeetham Amritapuri".

Query: Amrita Vishwa Vidyapeetham Amritapuri

key1

Maximum match, highest attention

Extract Values based on highest attention

The screenshot shows a YouTube search results page with the query "Amrita Vishwa Vidyapeetham Amritapuri" entered in the search bar. A red bracket labeled "key1" points from the left towards the first video result. A red arrow points from the text "Maximum match, highest attention" to the same video thumbnail. The video thumbnail is for a video titled "Amrita Vishwa Vidyapeetham, Amritapuri Campus, Kollam, Kerala" uploaded by "Amrita Vishwa Vidyapeetham" 2 years ago, with 152K views. The video duration is 20:32. The text "Extract Values based on highest attention" is displayed below the video thumbnail.

# Steps in Self Attention –

## First step

Create 3 vectors

- **Query vector (  $q_i$ )**
- **Key vector (  $k_i$ )**
- **Value vector (  $v_i$ )**

From each of the encoder's input vectors –ie the embedding of each word).

These 3 vectors are created by multiplying the embedding by three matrices that we trained during the training process

Input

Embedding

Queries

Keys

Values

Thinking

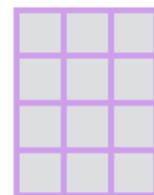
Machines

$x_1$

$x_2$

$q_1$

$q_2$



$W^Q$

$k_1$

$k_2$



$W^K$

$v_1$

$v_2$



$W^V$

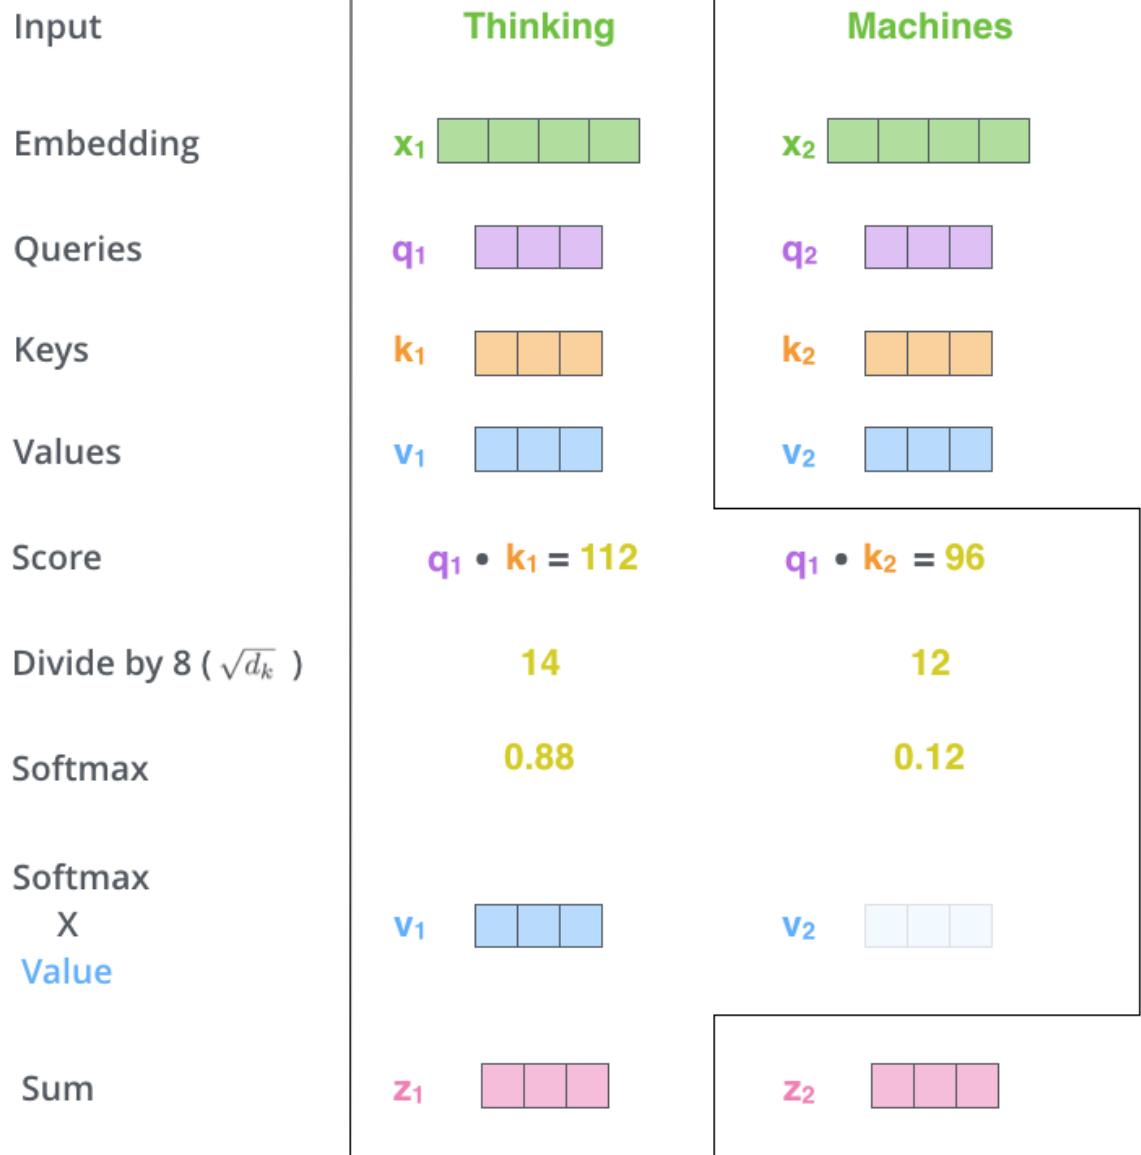
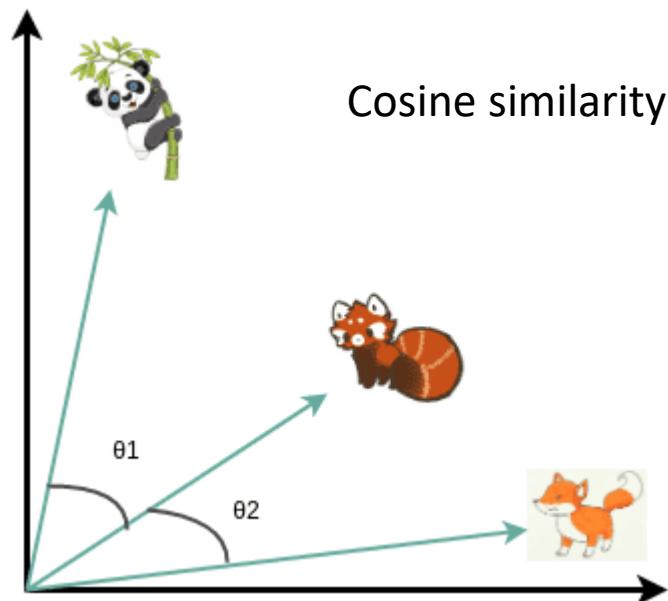
# Steps in Self Attention

## The second step

- Calculate a score : dot product of the **query vector** with the **key vector** of the respective word

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$



# Steps in Self Attention

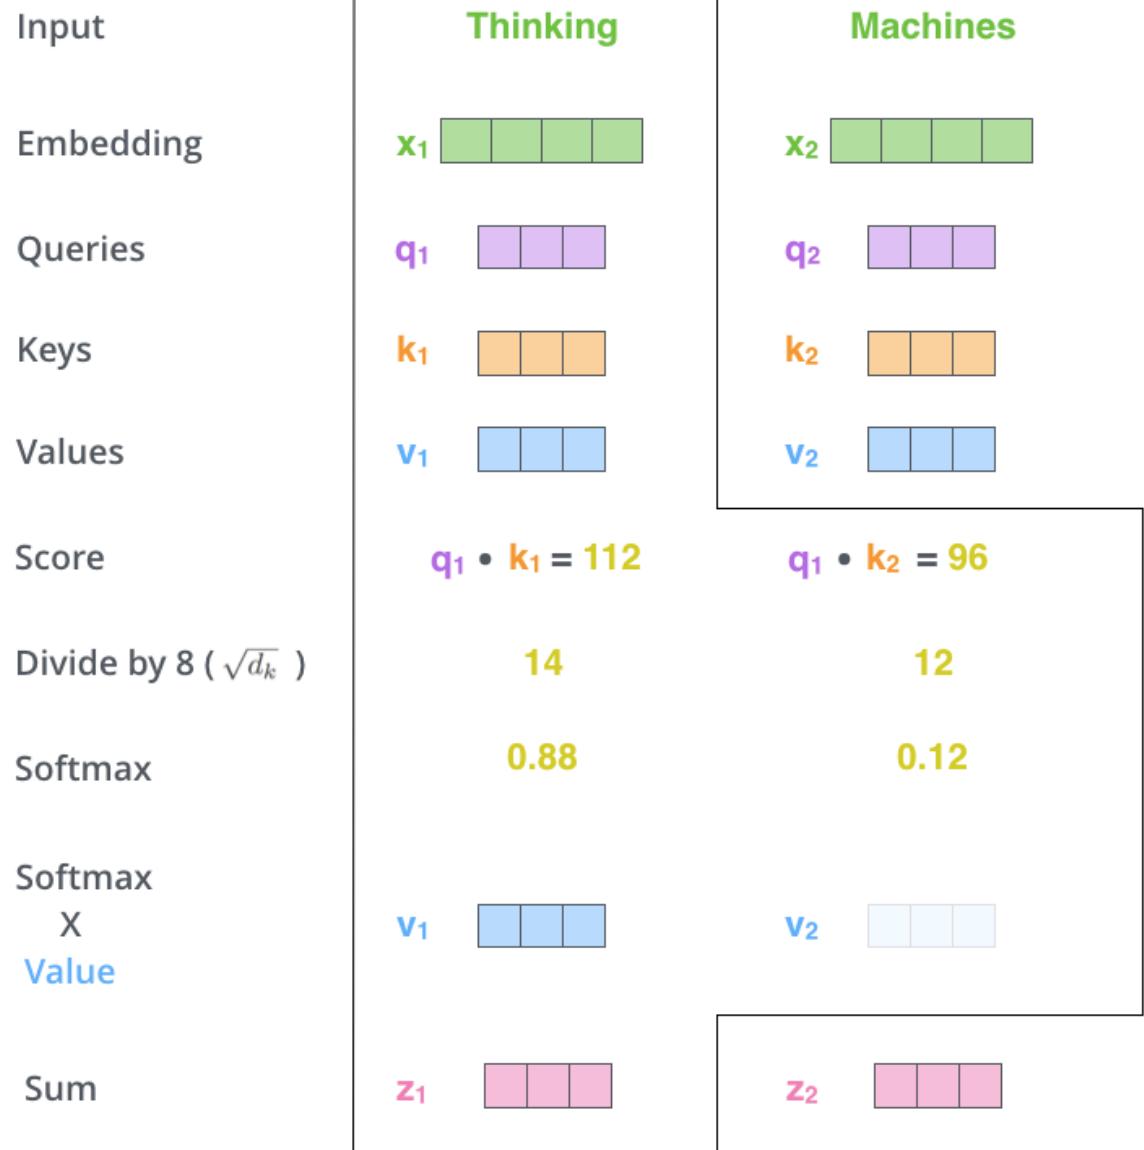
## Third step

Word	q vector	k vector	v vector	score	score / 8
thinking	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$
Machines		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$

Divide by 8 ( $\sqrt{d_k}$ )

Normalizing : This leads to having more stable gradients

The square root of the dimension of the key vectors used in the paper – 64.



# Steps in Self Attention –

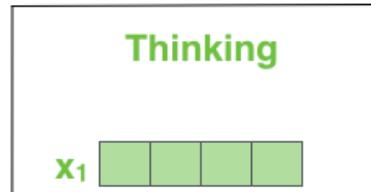
## Fourth step

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum
thinking	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$	$x_{11}$	$x_{11} * v_1$	$z_1$
Machines		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$	$x_{12}$	$x_{12} * v_2$	

Multiply each value vector by the softmax score

The intuition here is to keep intact the values of the word(s) we want to focus on, and drown-out irrelevant words

Input



Embedding



Queries



Keys



Values

$$q_1 \cdot k_1 = 112$$



Score

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$

Divide by 8 ( $\sqrt{d_k}$ )

$$14$$

$$12$$

Softmax

$$0.88$$

$$0.12$$

Softmax



X  
Value

Sum



# Steps in Self Attention –

## Fifth step

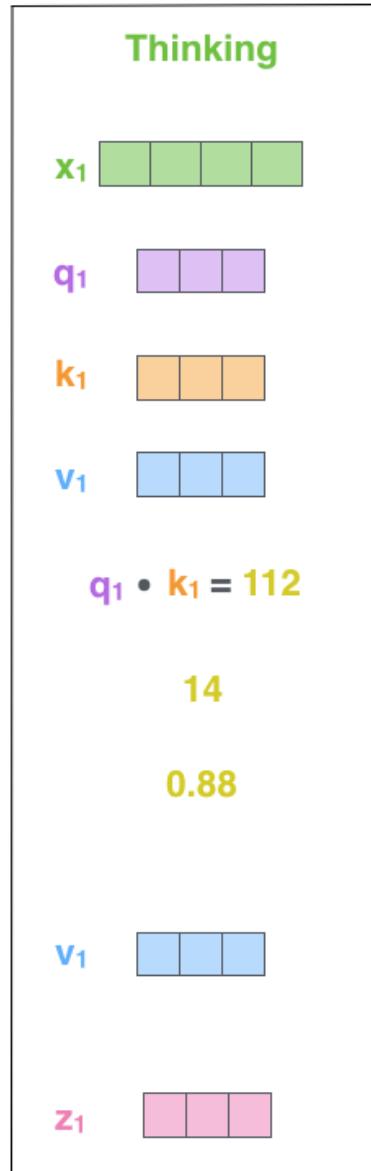
Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum
thinking	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$	$x_{11}$	$x_{11} * v_1$	$z_1$
Machines		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$	$x_{12}$	$x_{12} * v_2$	

Pass the result through a softmax operation.

Softmax normalizes the scores so they're all positive and add up to 1.

Sum up the weighted value vectors. This produces the output of the self-attention layer at this position (for the first word).

Input



Embedding

Queries

Keys

Values

Score

Divide by 8 ( $\sqrt{d_k}$ )

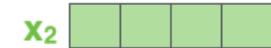
Softmax

Softmax

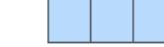
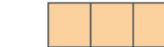
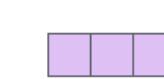
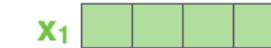
X  
Value

Sum

Machines



Thinking



$$q_1 \cdot k_1 = 112$$

$$14$$

$$0.88$$



$$q_1 \cdot k_2 = 96$$

$$12$$

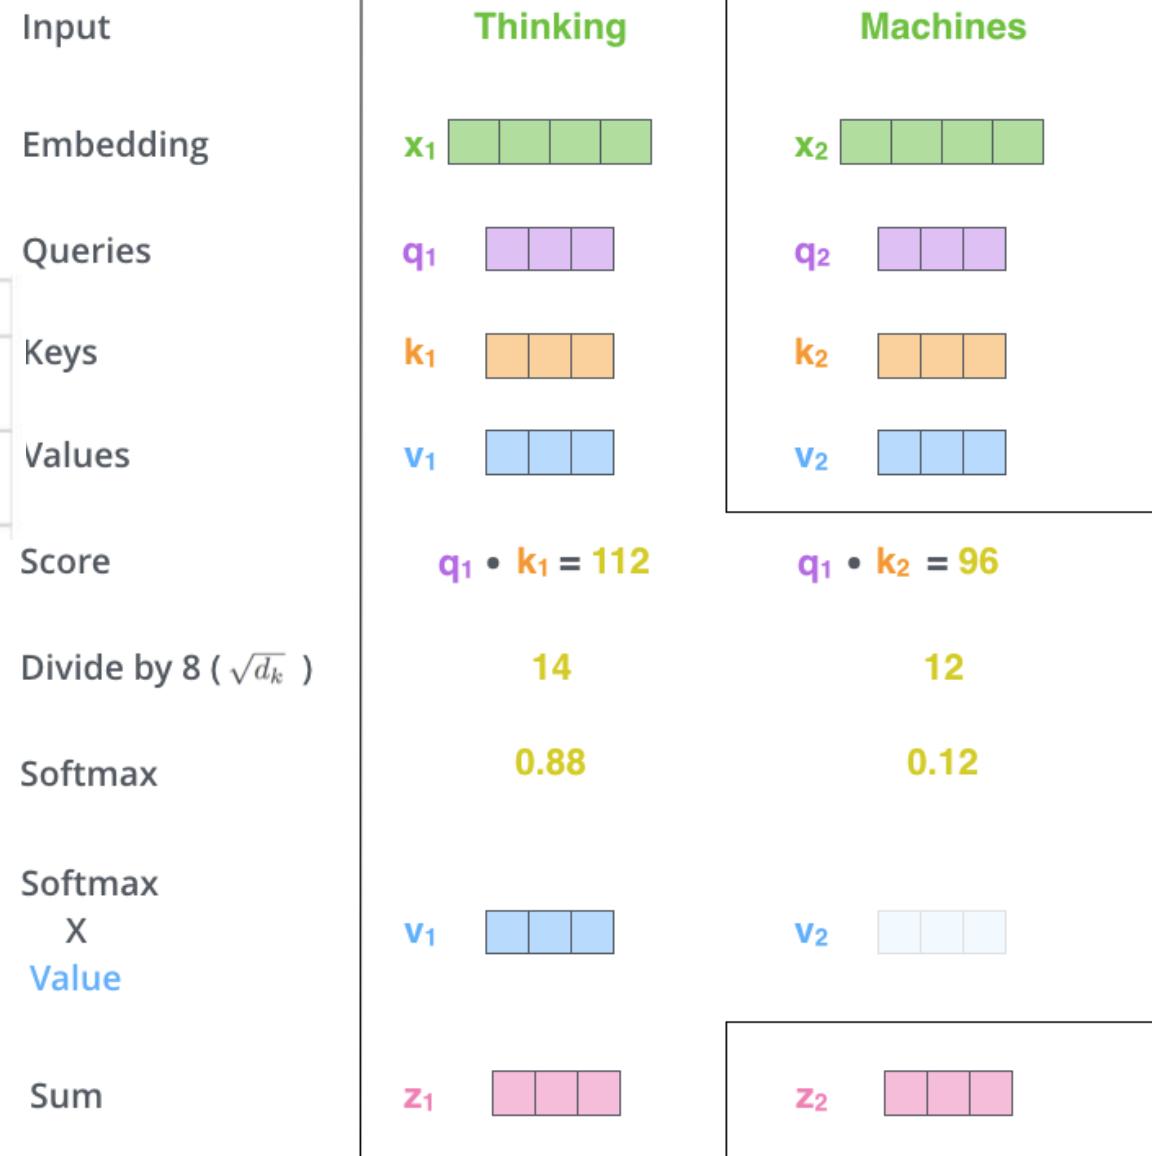
$$0.12$$



# Steps in Self Attention –

## Sixth Step

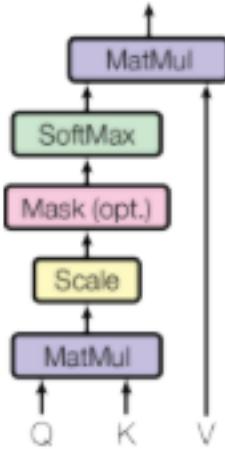
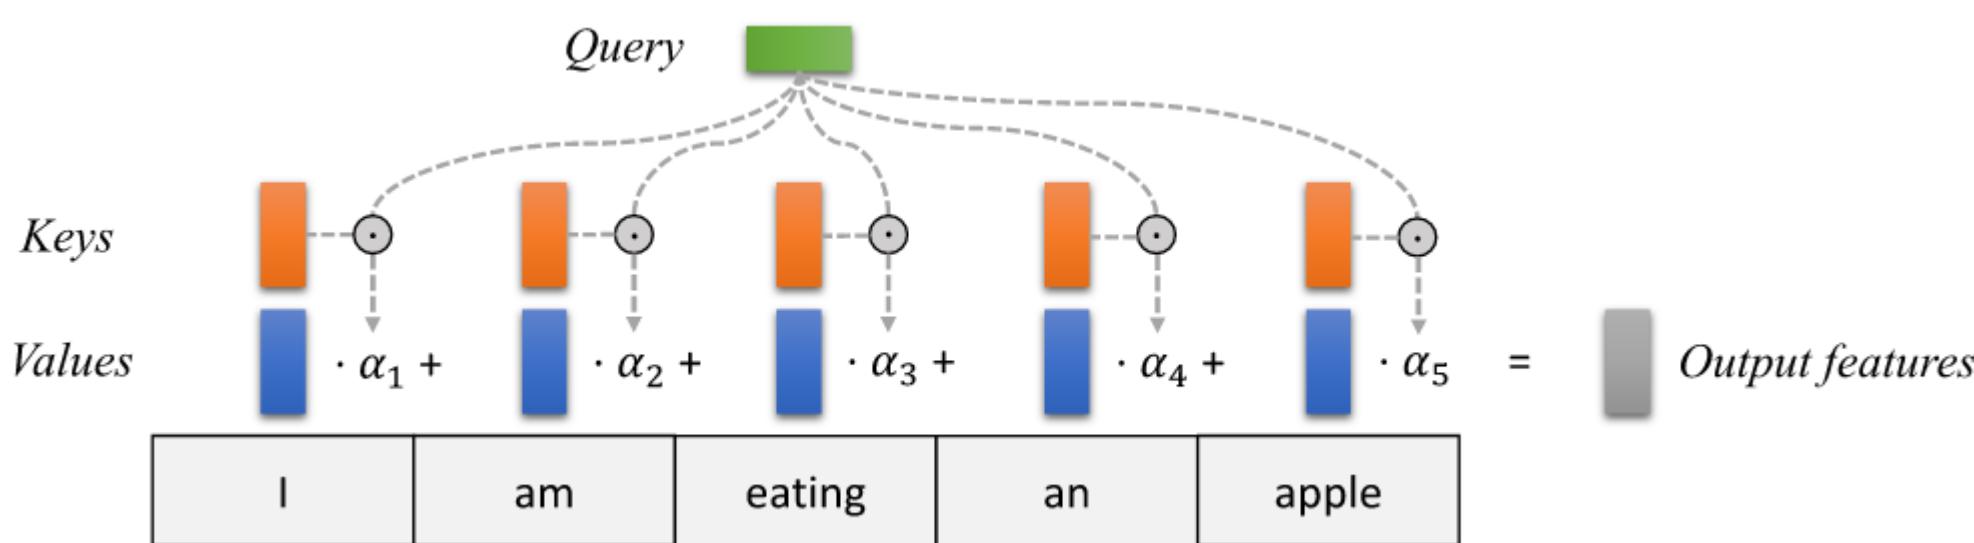
Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum <sup>#</sup>
thinking		$k_1$	$v_1$	$q_2 \cdot k_1$	$q_2 \cdot k_1 / 8$	$x_{21}$	$x_{21} * v_1$	
Machines	$q_2$	$k_2$	$v_2$	$q_2 \cdot k_2$	$q_2 \cdot k_2 / 8$	$x_{22}$	$x_{22} * v_2$	$z_2$



We need to score each word of the input sentence against this word. The score determines how much focus to place on other parts of the input sentence as we encode a word at a certain position.

# Self Attention –

Scaled Dot-Product Attention



$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

$$\alpha_i = \frac{\exp(f_{attn}(\text{key}_i, \text{query}))}{\sum_j \exp(f_{attn}(\text{key}_j, \text{query}))}, \quad \text{out} = \sum_i \alpha_i \cdot \text{value}_i$$

# Matrix Calculation of Self-Attention

$$\begin{array}{ccc} X & \times & W^Q \\ \begin{matrix} \textcolor{lightgreen}{X} \\ \hline \textcolor{lightgreen}{X} \end{matrix} & \times & \begin{matrix} \textcolor{violet}{W^Q} \\ \hline \textcolor{lightgray}{W^Q} \end{matrix} \\ & = & \begin{matrix} \textcolor{violet}{Q} \\ \hline \textcolor{lightgray}{Q} \end{matrix} \end{array}$$
$$\begin{array}{ccc} X & \times & W^K \\ \begin{matrix} \textcolor{lightgreen}{X} \\ \hline \textcolor{lightgreen}{X} \end{matrix} & \times & \begin{matrix} \textcolor{orange}{W^K} \\ \hline \textcolor{lightgray}{W^K} \end{matrix} \\ & = & \begin{matrix} \textcolor{orange}{K} \\ \hline \textcolor{lightgray}{K} \end{matrix} \end{array}$$
$$\begin{array}{ccc} X & \times & W^V \\ \begin{matrix} \textcolor{lightgreen}{X} \\ \hline \textcolor{lightgreen}{X} \end{matrix} & \times & \begin{matrix} \textcolor{blue}{W^V} \\ \hline \textcolor{lightgray}{W^V} \end{matrix} \\ & = & \begin{matrix} \textcolor{blue}{V} \\ \hline \textcolor{lightgray}{V} \end{matrix} \end{array}$$
$$\text{softmax}\left(\frac{\begin{matrix} Q & \times & K^T \\ \hline \textcolor{lightgray}{Q} & \times & \textcolor{lightgray}{K^T} \end{matrix}}{\sqrt{d_k}}\right) = Z$$

The diagram illustrates the matrix calculation of self-attention. It shows three parallel matrix multiplications of input  $X$  with weight matrices  $W^Q$ ,  $W^K$ , and  $W^V$  to produce  $Q$ ,  $K$ , and  $V$  respectively.  $Q$  and  $K$  are then multiplied together and divided by  $\sqrt{d_k}$  before being passed through a softmax function to produce  $Z$ .

# “Multi-headed” attention

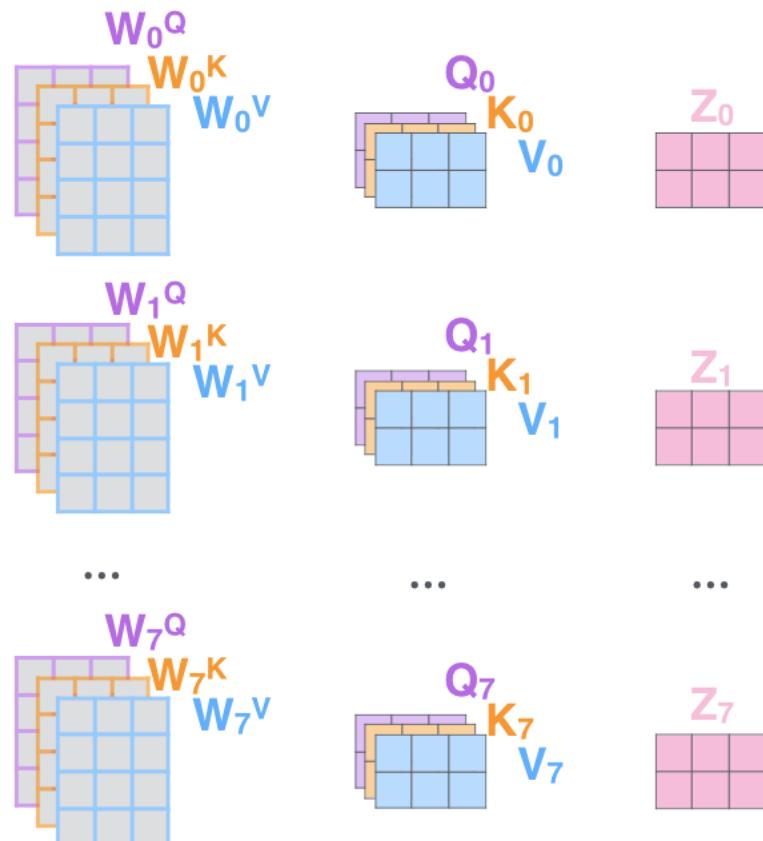


1) This is our input sentence\*    2) We embed each word\*

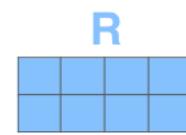
3) Split into 8 heads.  
We multiply X or R with weight matrices

4) Calculate attention using the resulting Q/K/V matrices

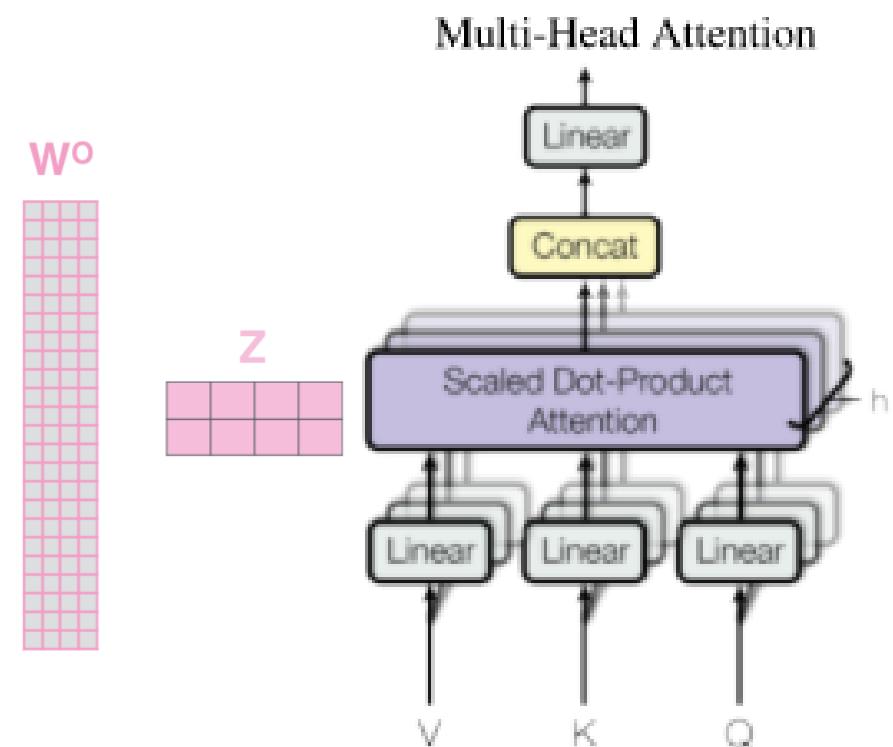
5) Concatenate the resulting Z matrices, then multiply with weight matrix W<sup>O</sup> to produce the output of the layer



\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



multiple “representation subspaces”



multiple sets of Query/Key/Value weight matrices

# Self Attention- Numerical Example

- The embedding of the word *I* is  $x_1 = [1.76, 2.22, \dots, 6.66]$ .
- The embedding of the word *am* is  $x_2 = [7.77, 0.631, \dots, 5.35]$ .
- The embedding of the word *good* is  $x_3 = [11.44, 10.10, \dots, 3.33]$ .

I	1.76	2.22	...	6.66	$\boldsymbol{x}_1$
am	7.77	0.631	...	5.35	$\boldsymbol{x}_2$
good	11.44	10.10	...	3.33	$\boldsymbol{x}_3$
					3x512
					X
input matrix (embedding matrix)					

Figure 1.6 – Input matrix

# Get Q K V

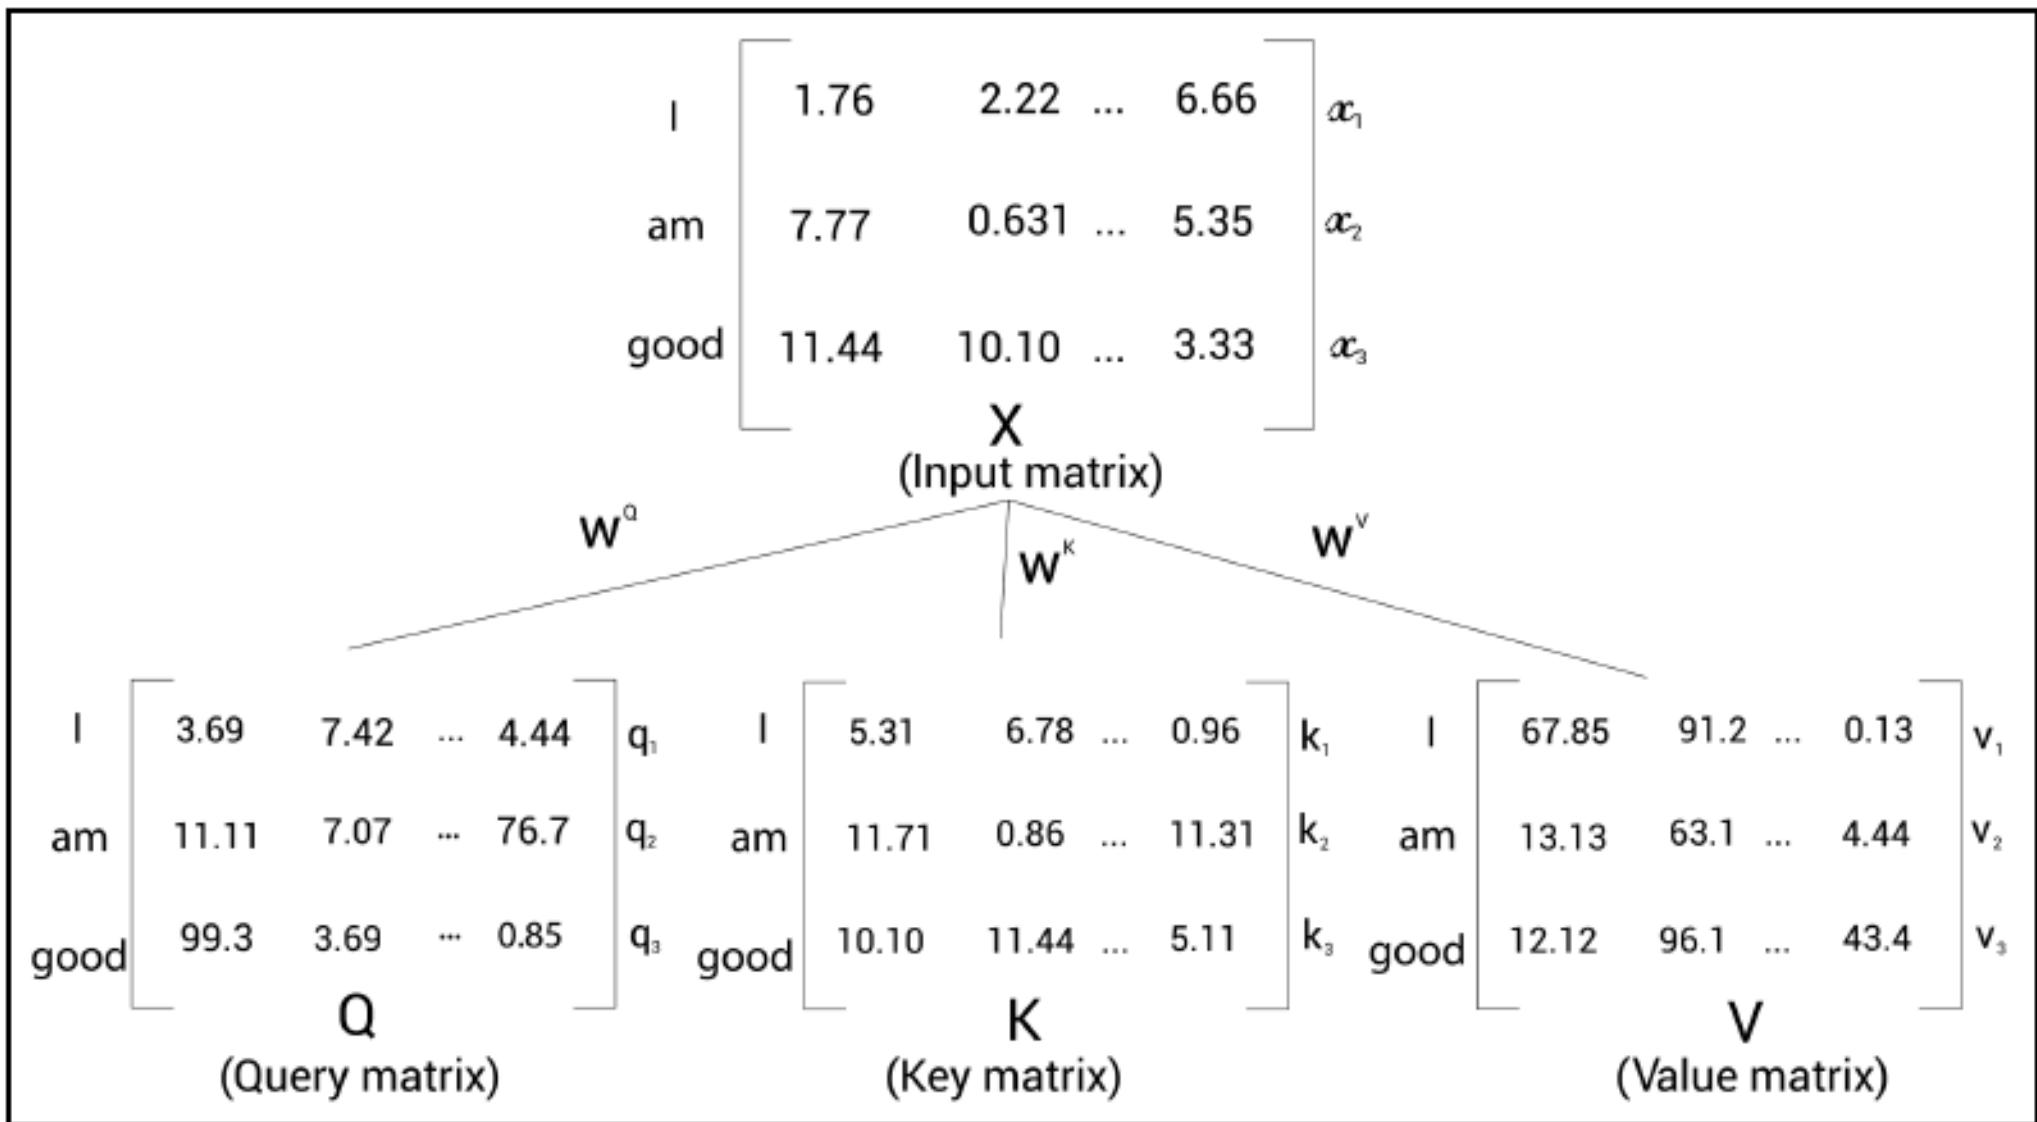


Figure 1.7 – Creating query, key, and value matrices

## Step 1

The first step in the self-attention mechanism is to compute the dot product between the query matrix,  $Q$ , and the key matrix,  $K^T$ :

$$\begin{matrix} & \begin{matrix} | & 3.69 & 7.42 \dots & 4.44 & | \\ am & 11.11 & 7.07 \dots & 76.7 & | \\ good & 99.3 & 3.69 \dots & 0.85 & | \end{matrix} & \begin{matrix} q_1 & | & 5.31 & 11.71 & 10.10 \\ q_2 & | & 6.78 & 0.86 & 11.44 \\ q_3 & | & \vdots & \vdots & \vdots \\ k_1 & | & 0.96 & 11.31 & 5.11 \\ k_2 & | & & & \\ k_3 & | & & & \end{matrix} \\ Q & & K^T \end{matrix}$$

Figure 1.9 – Query and key matrices

## Step 2

The next step in the self-attention mechanism is to divide the  $Q \cdot K^T$  matrix by the square root of the dimension of the key vector. But why do we have to do that? This is useful in obtaining stable gradients.

$$\begin{array}{c} \left| \begin{array}{c|ccc} & 3.69 & 7.42 & \dots & 4.44 \\ am & 11.11 & 7.07 & \dots & 76.7 \\ good & 99.3 & 3.69 & \dots & 0.85 \end{array} \right| q_1 \quad \left| \begin{array}{c|ccc} & 5.31 & 11.71 & 10.10 \\ am & 6.78 & 0.86 & 11.44 \\ good & 0.96 & 11.31 & 5.11 \end{array} \right| k_1 \\ Q \qquad \qquad \qquad K^T \end{array} = \begin{array}{c} \left| \begin{array}{c|ccc} & q_1 \cdot k_1 & q_1 \cdot k_2 & q_1 \cdot k_3 \\ am & q_2 \cdot k_1 & q_2 \cdot k_2 & q_2 \cdot k_3 \\ good & q_3 \cdot k_1 & q_3 \cdot k_2 & q_3 \cdot k_3 \end{array} \right| \\ \left| \begin{array}{c|ccc} & 110 & 90 & 80 \\ am & 70 & 99 & 70 \\ good & 90 & 70 & 100 \end{array} \right| \\ QK^T = \end{array}$$

Figure 1.10 – Computing the dot product between the query and key matrices

# Normalising the scores

$$\frac{QK^T}{\sqrt{d_k}} = \frac{QK^T}{8} = \begin{matrix} & | & am & good \\ I & | & \begin{bmatrix} 13.75 & 11.25 & 10 \\ 8.75 & 12.375 & 8.75 \\ 11.25 & 8.75 & 12.5 \end{bmatrix} \\ & | & am & good \end{matrix}$$

Figure 1.14 – Dividing  $QK^T$  by the square root of  $d_k$

$$\text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) = \begin{matrix} & | & am & good \\ I & | & \begin{bmatrix} 0.90 & 0.07 & 0.03 \\ 0.025 & 0.95 & 0.025 \\ 0.21 & 0.03 & 0.76 \end{bmatrix} \\ & | & am & good \end{matrix}$$

Figure 1.15 – Applying the softmax function

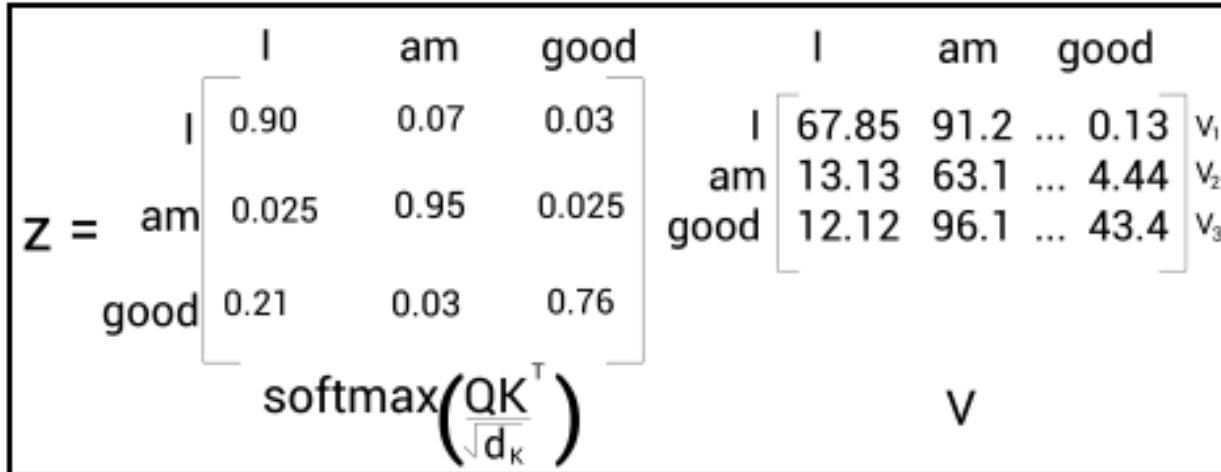


Figure 1.16 – Computing the attention matrix

$$Z = \begin{bmatrix} z_1 \\ z_2 \\ z_3 \end{bmatrix} \begin{array}{l} | \\ am \\ | \\ good \end{array}$$

Figure 1.17 – Result of the attention matrix

$$z_1 = 0.90 [67.85 \ 91.2 \ \dots] + 0.07 [13.13 \ 63.1 \ \dots] + 0.03 [12.12 \ 96.1 \ \dots]$$

$v_1(I)$                      $v_2(am)$                      $v_3(good)$

Figure 1.18 – Self-attention of the word I

$$Z_2 = 0.025 \begin{bmatrix} 67.85 \\ 91.2 \\ \dots \end{bmatrix} + 0.95 \begin{bmatrix} 13.13 \\ 63.1 \\ \dots \end{bmatrix} + 0.025 \begin{bmatrix} 12.12 \\ 96.1 \\ \dots \end{bmatrix}$$

$v_1(l)$                      $v_2(am)$                      $v_3(good)$

Figure 1.20 – Self-attention of the word am

$$Z_3 = 0.21 \begin{bmatrix} 67.85 \\ 91.2 \\ \dots \end{bmatrix} + 0.03 \begin{bmatrix} 13.13 \\ 63.1 \\ \dots \end{bmatrix} + 0.76 \begin{bmatrix} 12.12 \\ 96.1 \\ \dots \end{bmatrix}$$

$v_1(l)$                      $v_2(am)$                      $v_3(good)$

Figure 1.21 – Self-attention of the word good

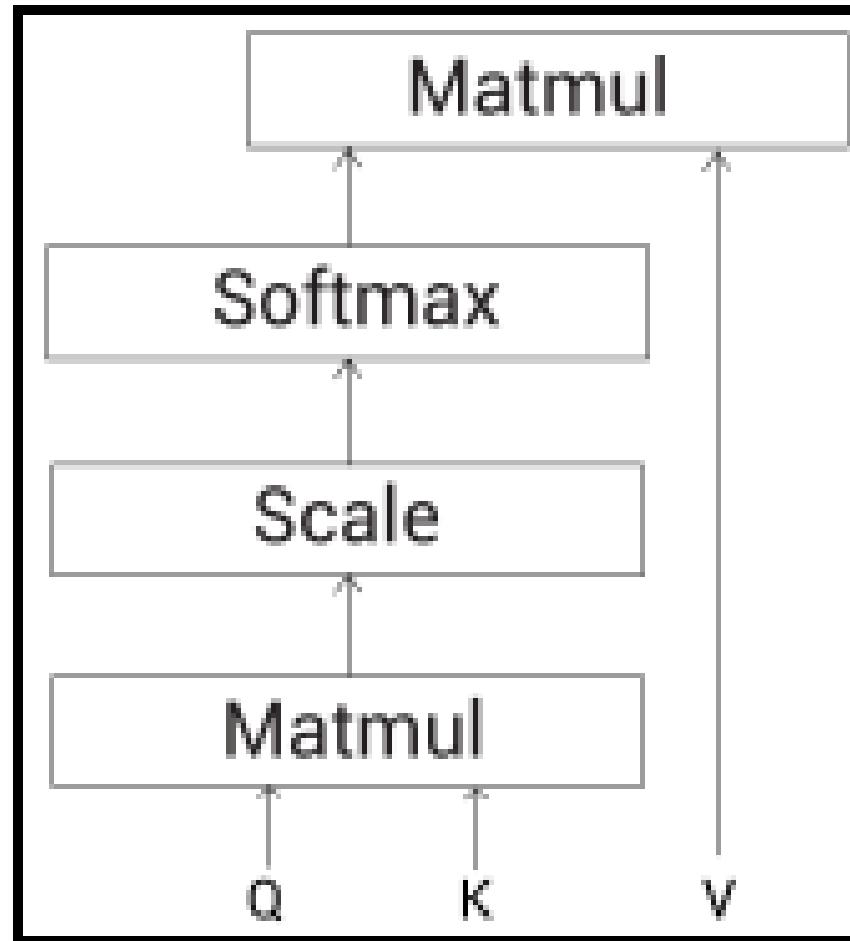
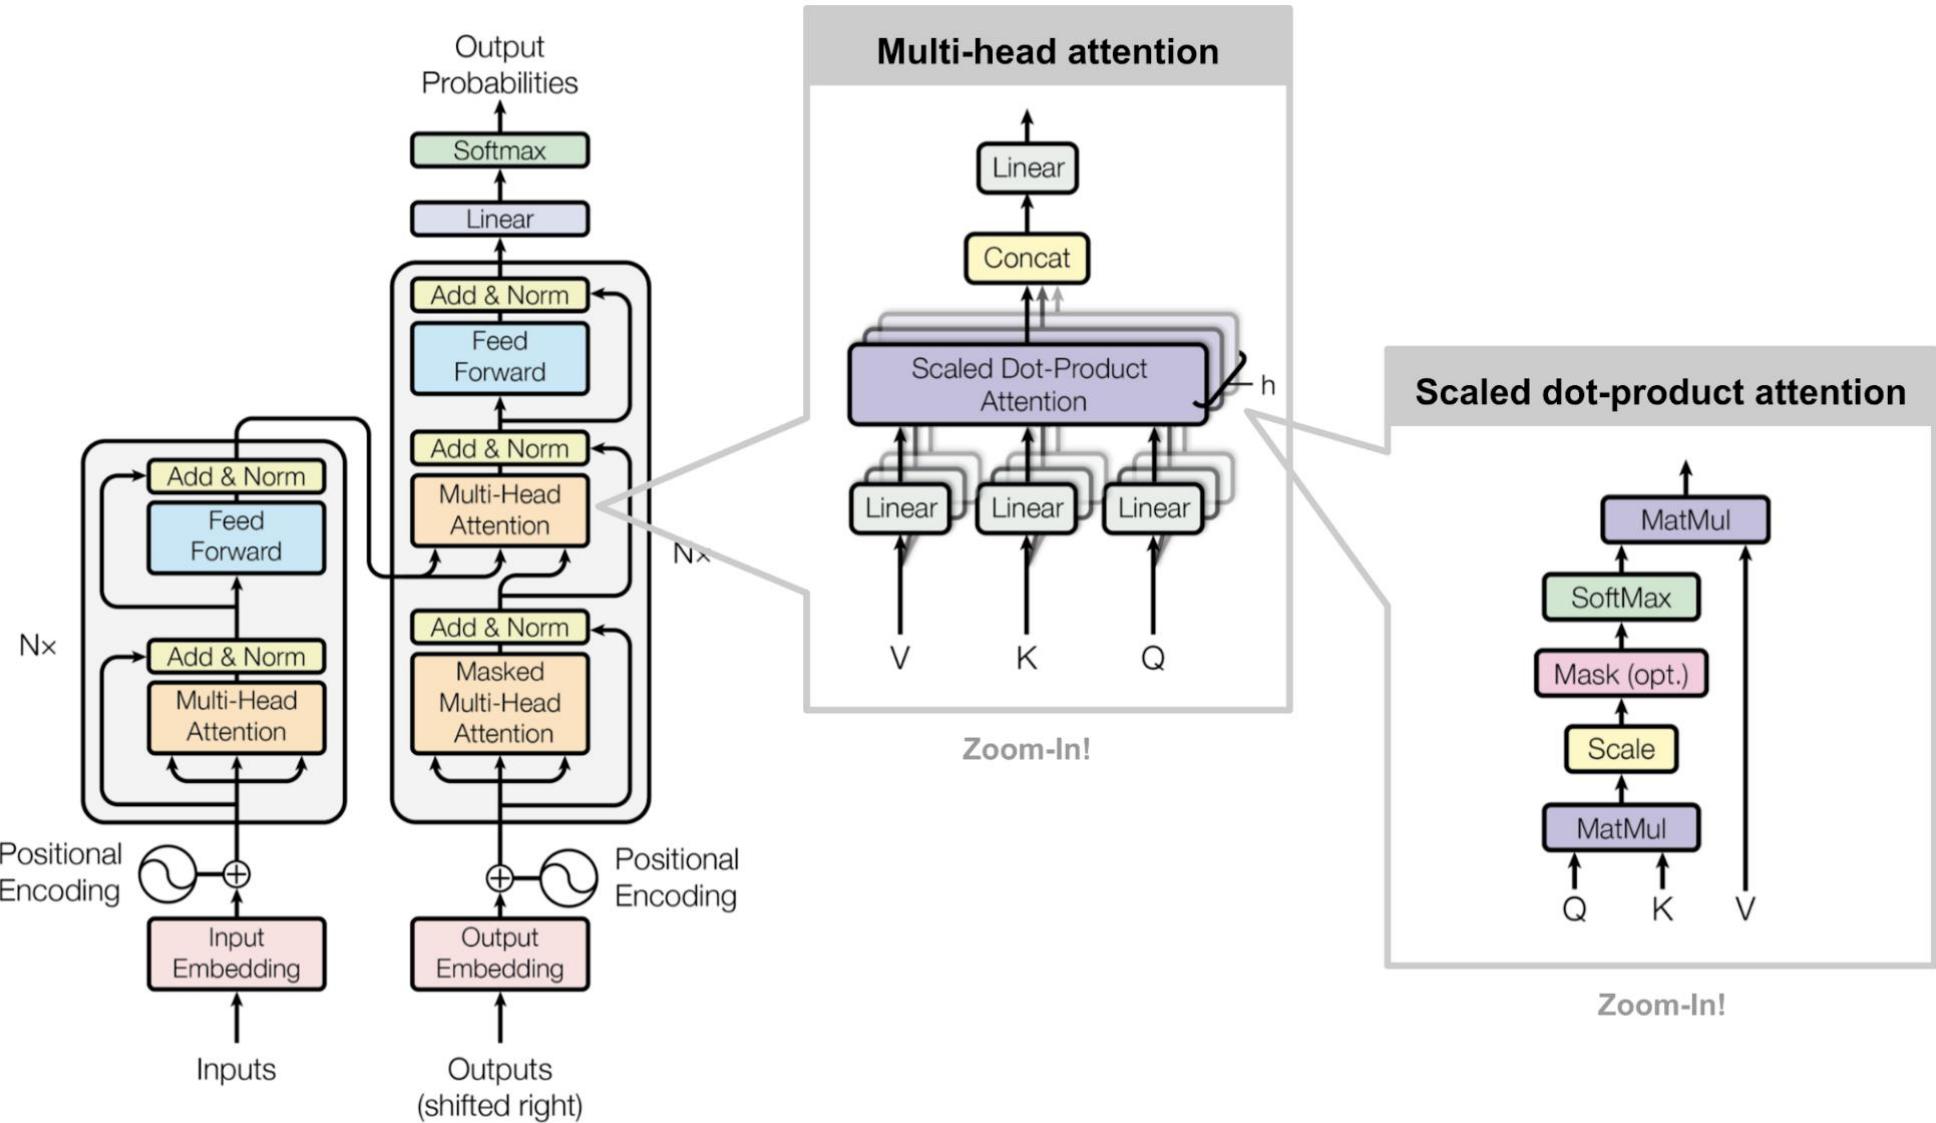
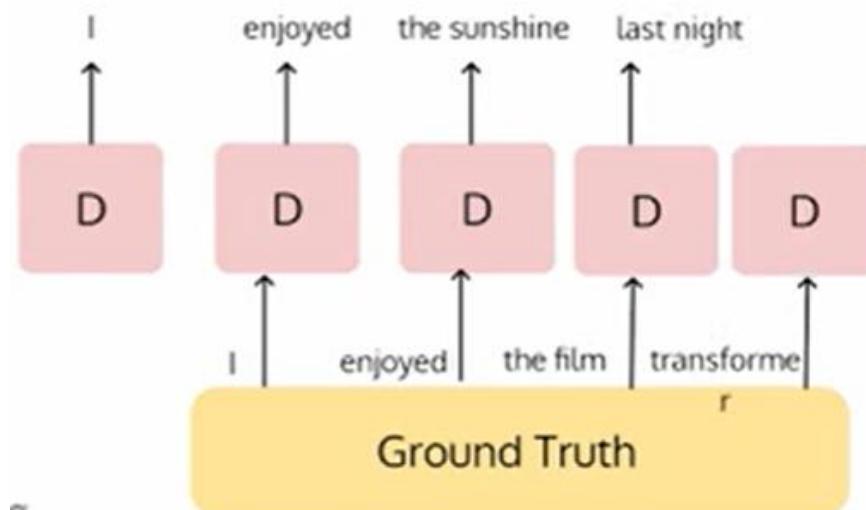
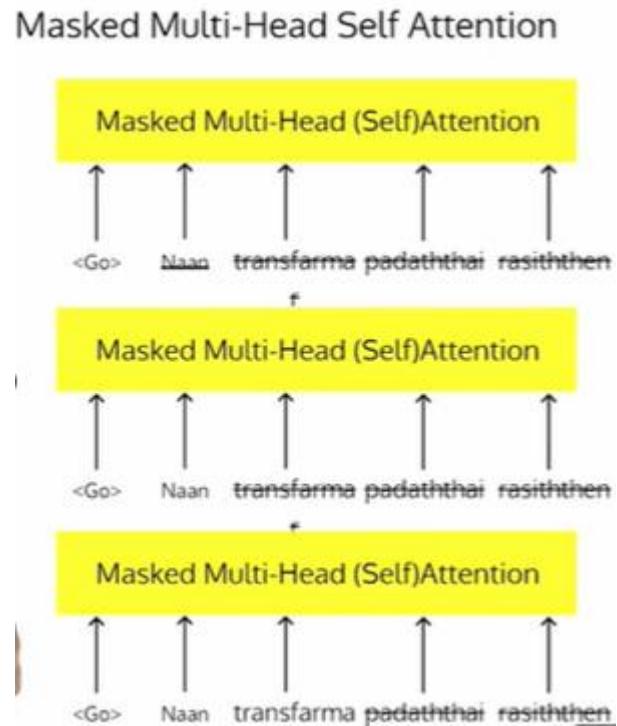


Figure 1.22 – Self-attention mechanism

# Transformers



# Masked Multi Head Attention



$$e^0 = 1$$
$$e^{-\infty} = 0$$

## Masking in Multi-Head Attention: Why Is It Needed?

- ✓ Prevents Attending to Unimportant Tokens (e.g., padding).
- ✓ Ensures Autoregressive Decoding (prevents future peeking).
- ✓ Maintains Training Consistency (aligns with left-to-right generation).
- ✓ Optimizes Computation (avoids unnecessary attention to empty inputs).

Masking is done by inserting negative infinite at the respective positions.

$q_1 \cdot k_1$	$-\infty$	$-\infty$	$-\infty$	$-\infty$
$q_2 \cdot k_1$	$q_2 \cdot k_2$	$-\infty$	$-\infty$	$-\infty$
$q_3 \cdot k_1$	$q_3 \cdot k_2$	$q_3 \cdot k_3$	$-\infty$	$-\infty$
$q_4 \cdot k_1$	$q_4 \cdot k_2$	$q_4 \cdot k_3$	$q_4 \cdot k_4$	$-\infty$
$q_5 \cdot k_1$	$q_5 \cdot k_2$	$q_5 \cdot k_3$	$q_5 \cdot k_4$	$q_5 \cdot k_5$

$T$

$T$

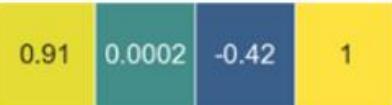
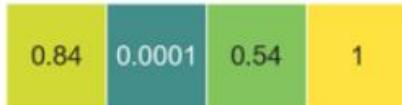
This actually forms an triangular matrix with one half all zeros and the other half all negative infinity

$$M = \begin{bmatrix} 0 & -\infty & -\infty & -\infty & -\infty \\ 0 & 0 & -\infty & -\infty & -\infty \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

# Representing The Order of The Sequence Using Positional Encoding

Multi-Head Attention block is permutation-equivariant, w.r.t its inputs and cannot distinguish whether an input comes before another one in the sequence or not.- hence positional encoding

Positional  
encoding

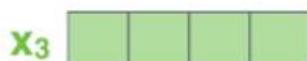


+

+

+

Embedding



Input

je

suis

étudiant

$$PE_{(pos,i)} = \begin{cases} \sin\left(\frac{pos}{10000^{i/d_{model}}}\right) & \text{if } i \bmod 2 = 0 \\ \cos\left(\frac{pos}{10000^{(i-1)/d_{model}}}\right) & \text{otherwise} \end{cases}$$

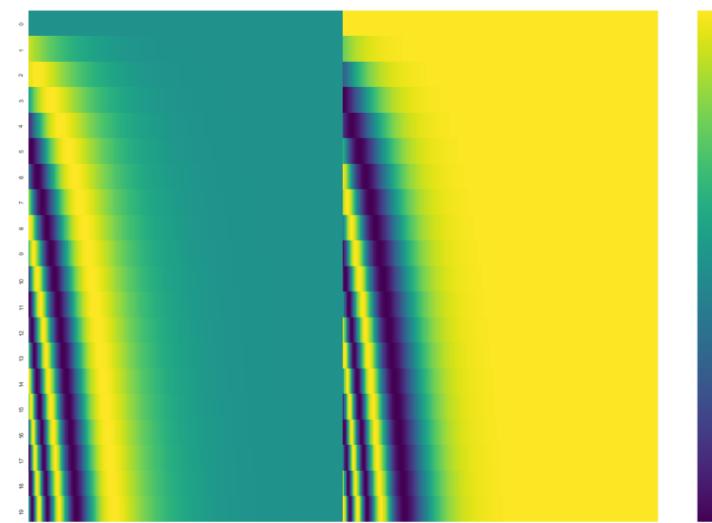
$PE_{(pos,i)}$

- pos is the position of the token in the sequence.
- i is the dimension index within the positional encoding.
- dmodel is the dimensionality of the word embedding .

Vaswani et al 2017

Permutation-equivariant, :means that if we switch two input elements in the sequence, e.g.  $X_1 \leftrightarrow X_2$  the output is exactly the same besides the elements 1 and 2 switched

A real example of positional encoding for 20 words (rows) with an embedding size of 512 (columns)

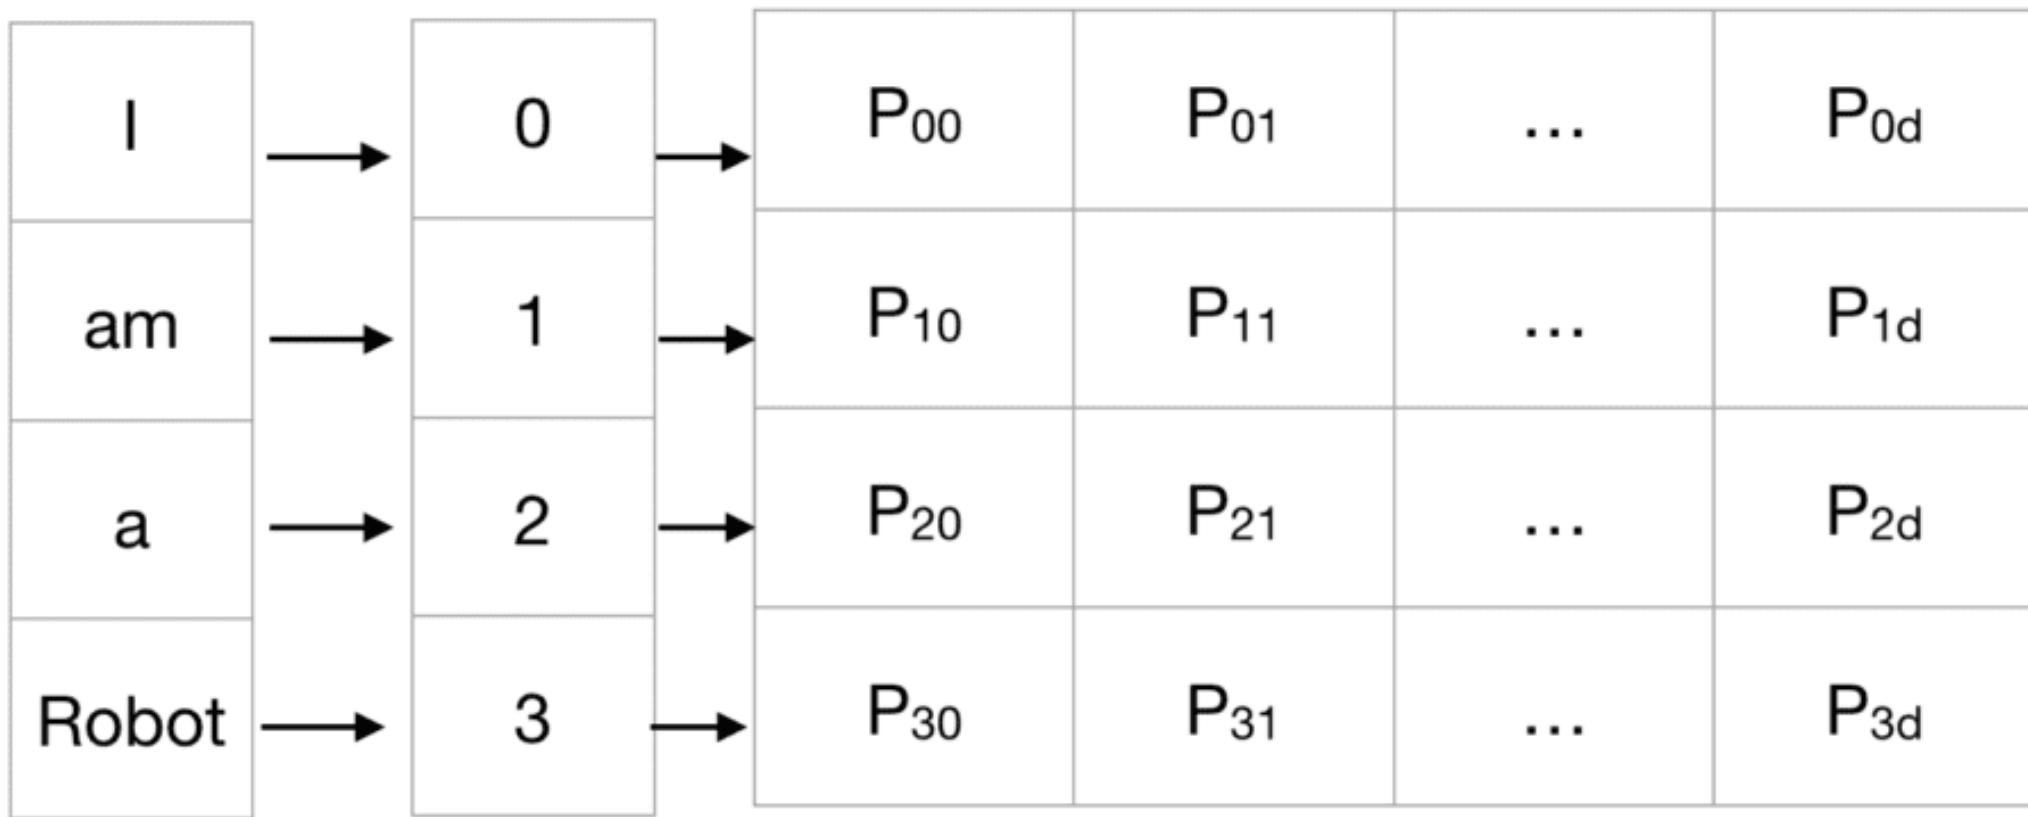


Sequence      Index  
of token

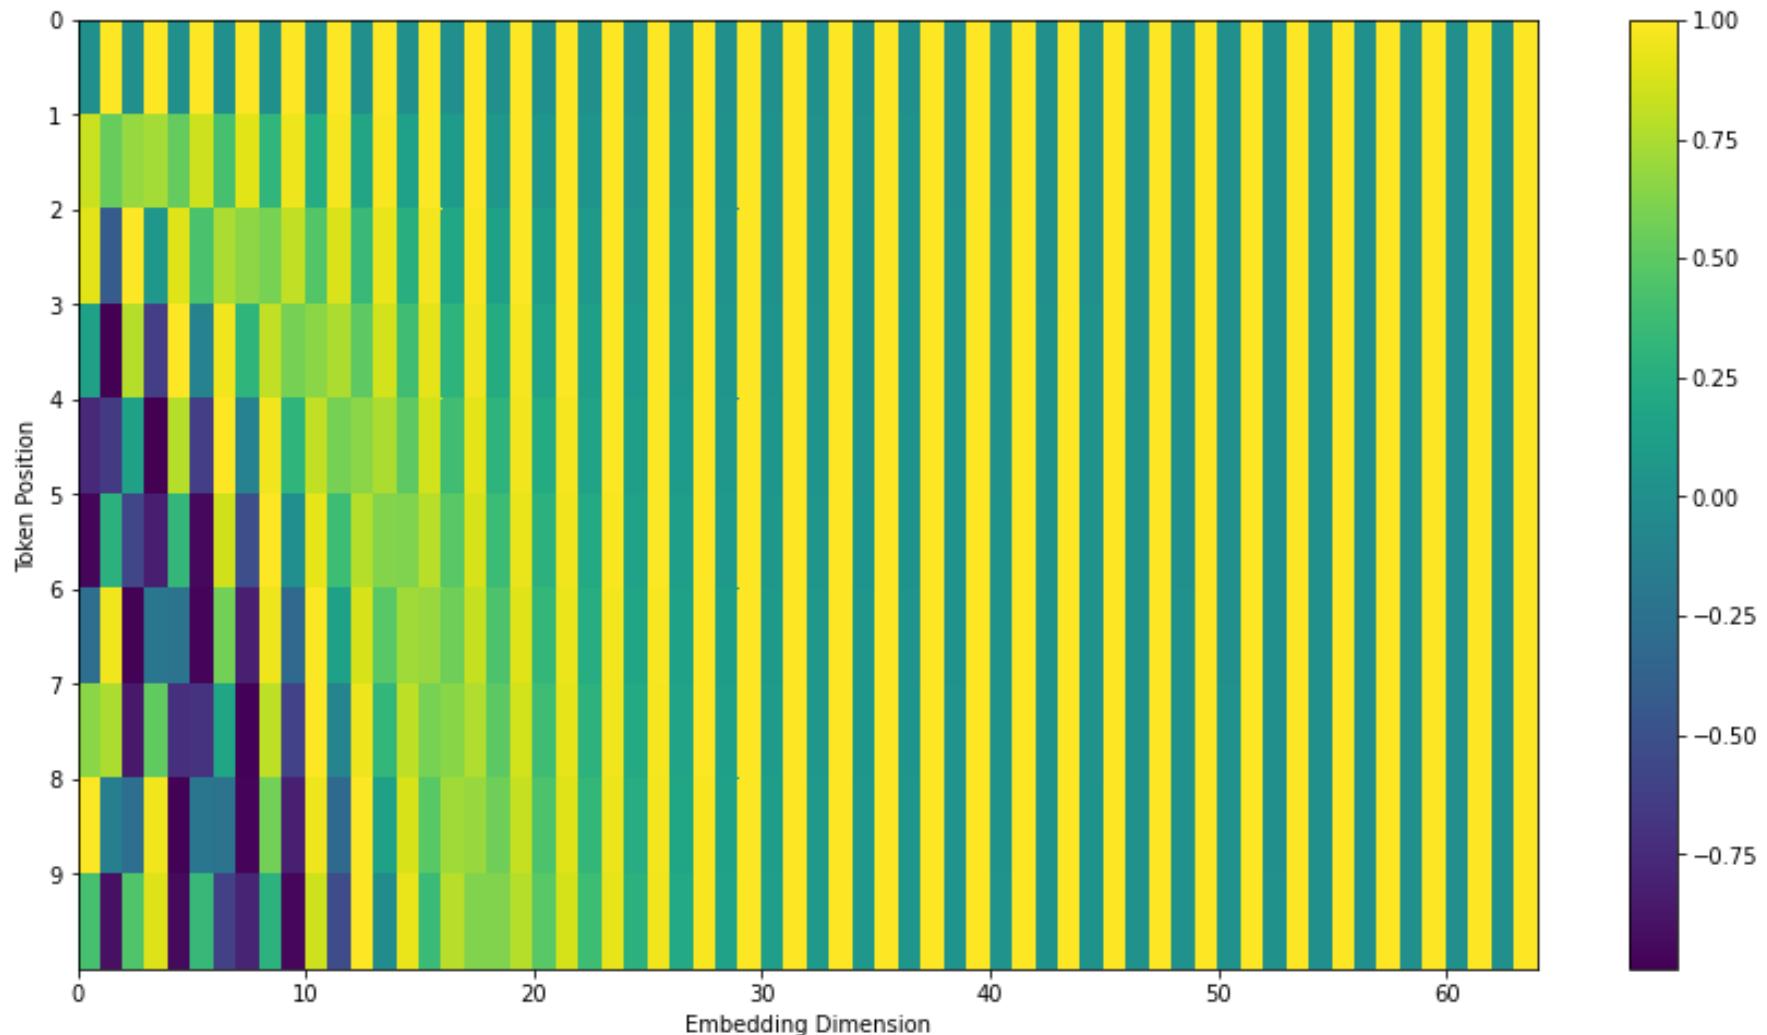
$$d_{\text{model}} = d+1$$

$$P_{pos,i}$$

Positional Encoding Matrix



Positional Encoding Matrix for the sequence 'I am a robot'



	1.769	2.22	3.4	5.8	$x_1$
X = am	7.3	9.9	8.5	7.1	$x_2$
good	9.1	7.1	0.85	10.1	$x_3$

Input matrix  
(embedding matrix)

$$P(\text{pos}, 2i) = \sin\left(\frac{\text{pos}}{1000^{2i/d_{\text{model}}}}\right)$$

$$P(\text{pos}, 2i + 1) = \cos\left(\frac{\text{pos}}{1000^{2i/d_{\text{model}}}}\right)$$

$d_{\text{model}}$  is embedding dimension

$\text{pos}$  is the position of the word in the sentence

$i$  is the position of the embedding

$$d_{\text{model}}=4$$

$$\begin{aligned} X &= \begin{bmatrix} 1.769 & 2.22 & 3.4 & 5.8 \\ 7.3 & 9.9 & 8.5 & 7.1 \\ 9.1 & 7.1 & 0.85 & 10.1 \end{bmatrix} + \begin{bmatrix} 0 & 1 & 0 & 1 \\ 0.841 & 0.54 & 0.01 & 0.99 \\ 0.909 & -0.416 & 0.02 & 0.99 \end{bmatrix} \\ &\quad X \\ &= \begin{bmatrix} 1.769 & 3.22 & 3.4 & 6.8 \\ 8.14 & 10.44 & 8.51 & 8.09 \\ 10.0 & 6.68 & 0.87 & 11.09 \end{bmatrix} \end{aligned}$$

Add input and positional encoding matrix

$$P = \begin{matrix} & | \sin\left(\frac{\text{pos}}{10000^0}\right) \cos\left(\frac{\text{pos}}{10000^0}\right) \sin\left(\frac{\text{pos}}{10000^{2/4}}\right) \cos\left(\frac{\text{pos}}{10000^{2/4}}\right) \\ \text{am} & | \sin\left(\frac{\text{pos}}{10000^0}\right) \cos\left(\frac{\text{pos}}{10000^0}\right) \sin\left(\frac{\text{pos}}{10000^{2/4}}\right) \cos\left(\frac{\text{pos}}{10000^{2/4}}\right) \\ \text{good} & | \sin\left(\frac{\text{pos}}{10000^0}\right) \cos\left(\frac{\text{pos}}{10000^0}\right) \sin\left(\frac{\text{pos}}{10000^{2/4}}\right) \cos\left(\frac{\text{pos}}{10000^{2/4}}\right) \end{matrix}$$

Figure 1.27 – Computing the positional encoding matrix

$$P = \begin{matrix} & | \sin(\text{pos}) \cos(\text{pos}) \sin\left(\frac{\text{pos}}{100}\right) \cos\left(\frac{\text{pos}}{100}\right) \\ \text{am} & | \sin(\text{pos}) \cos(\text{pos}) \sin\left(\frac{\text{pos}}{100}\right) \cos\left(\frac{\text{pos}}{100}\right) \\ \text{good} & | \sin(\text{pos}) \cos(\text{pos}) \sin\left(\frac{\text{pos}}{100}\right) \cos\left(\frac{\text{pos}}{100}\right) \end{matrix}$$

Figure 1.28 – Computing the positional encoding matrix

$$P = \begin{matrix} & | \sin(0) \cos(0) \sin(0/100) \cos(0/100) \\ \text{am} & | \sin(1) \cos(1) \sin(1/100) \cos(1/100) \\ \text{good} & | \sin(2) \cos(2) \sin(2/100) \cos(2/100) \end{matrix}$$

Figure 1.29 – Computing the positional encoding matrix

$$P = \begin{matrix} & | 0 \quad 1 \quad 0 \quad 1 \\ \text{am} & | 0.841 \quad 0.540 \quad 0.009 \quad 0.999 \\ \text{good} & | 0.909 \quad -0.416 \quad 0.019 \quad 0.999 \end{matrix}$$

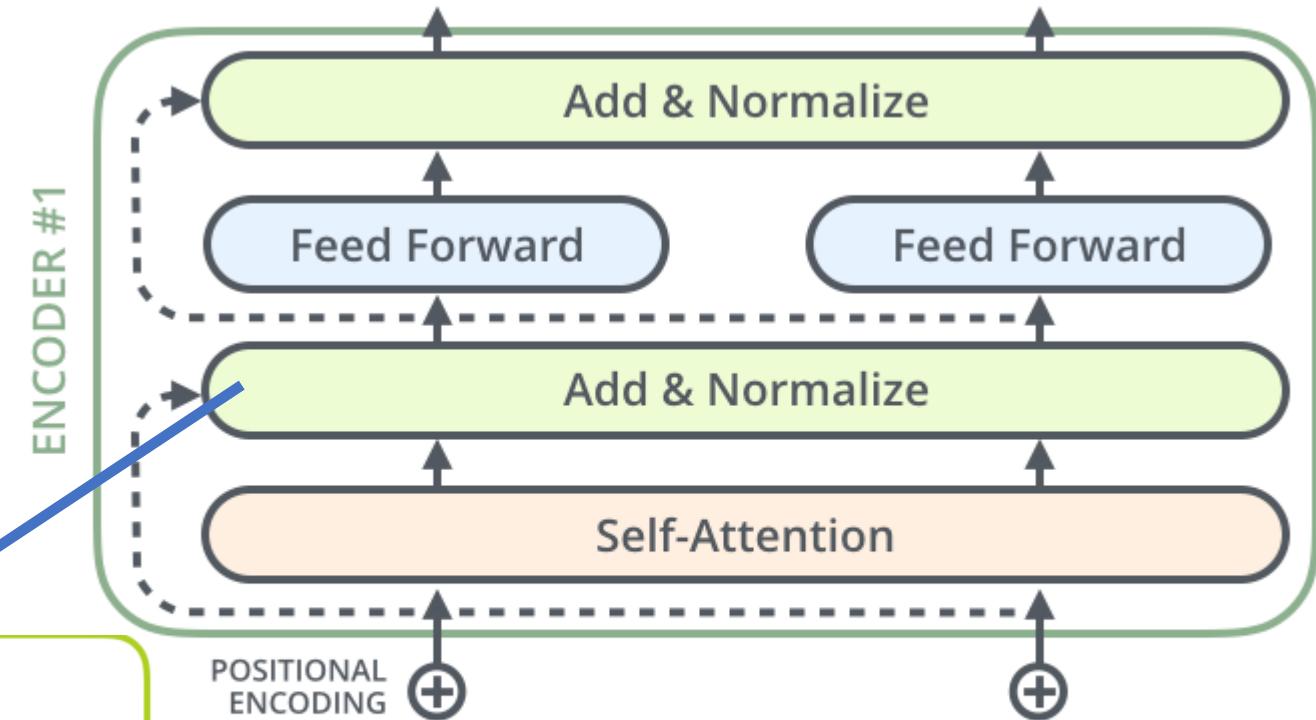
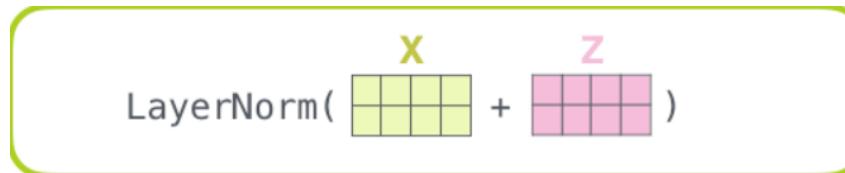
Figure 1.30 – Positional encoding matrix

# Residual connection, Layer Norm, FFN

Each encoder has a

- Residual connection around it,
- Followed by a layer normalization step
- Feed forward : -deepens our network, employing linear layers to analyse patterns in the attention layers output.

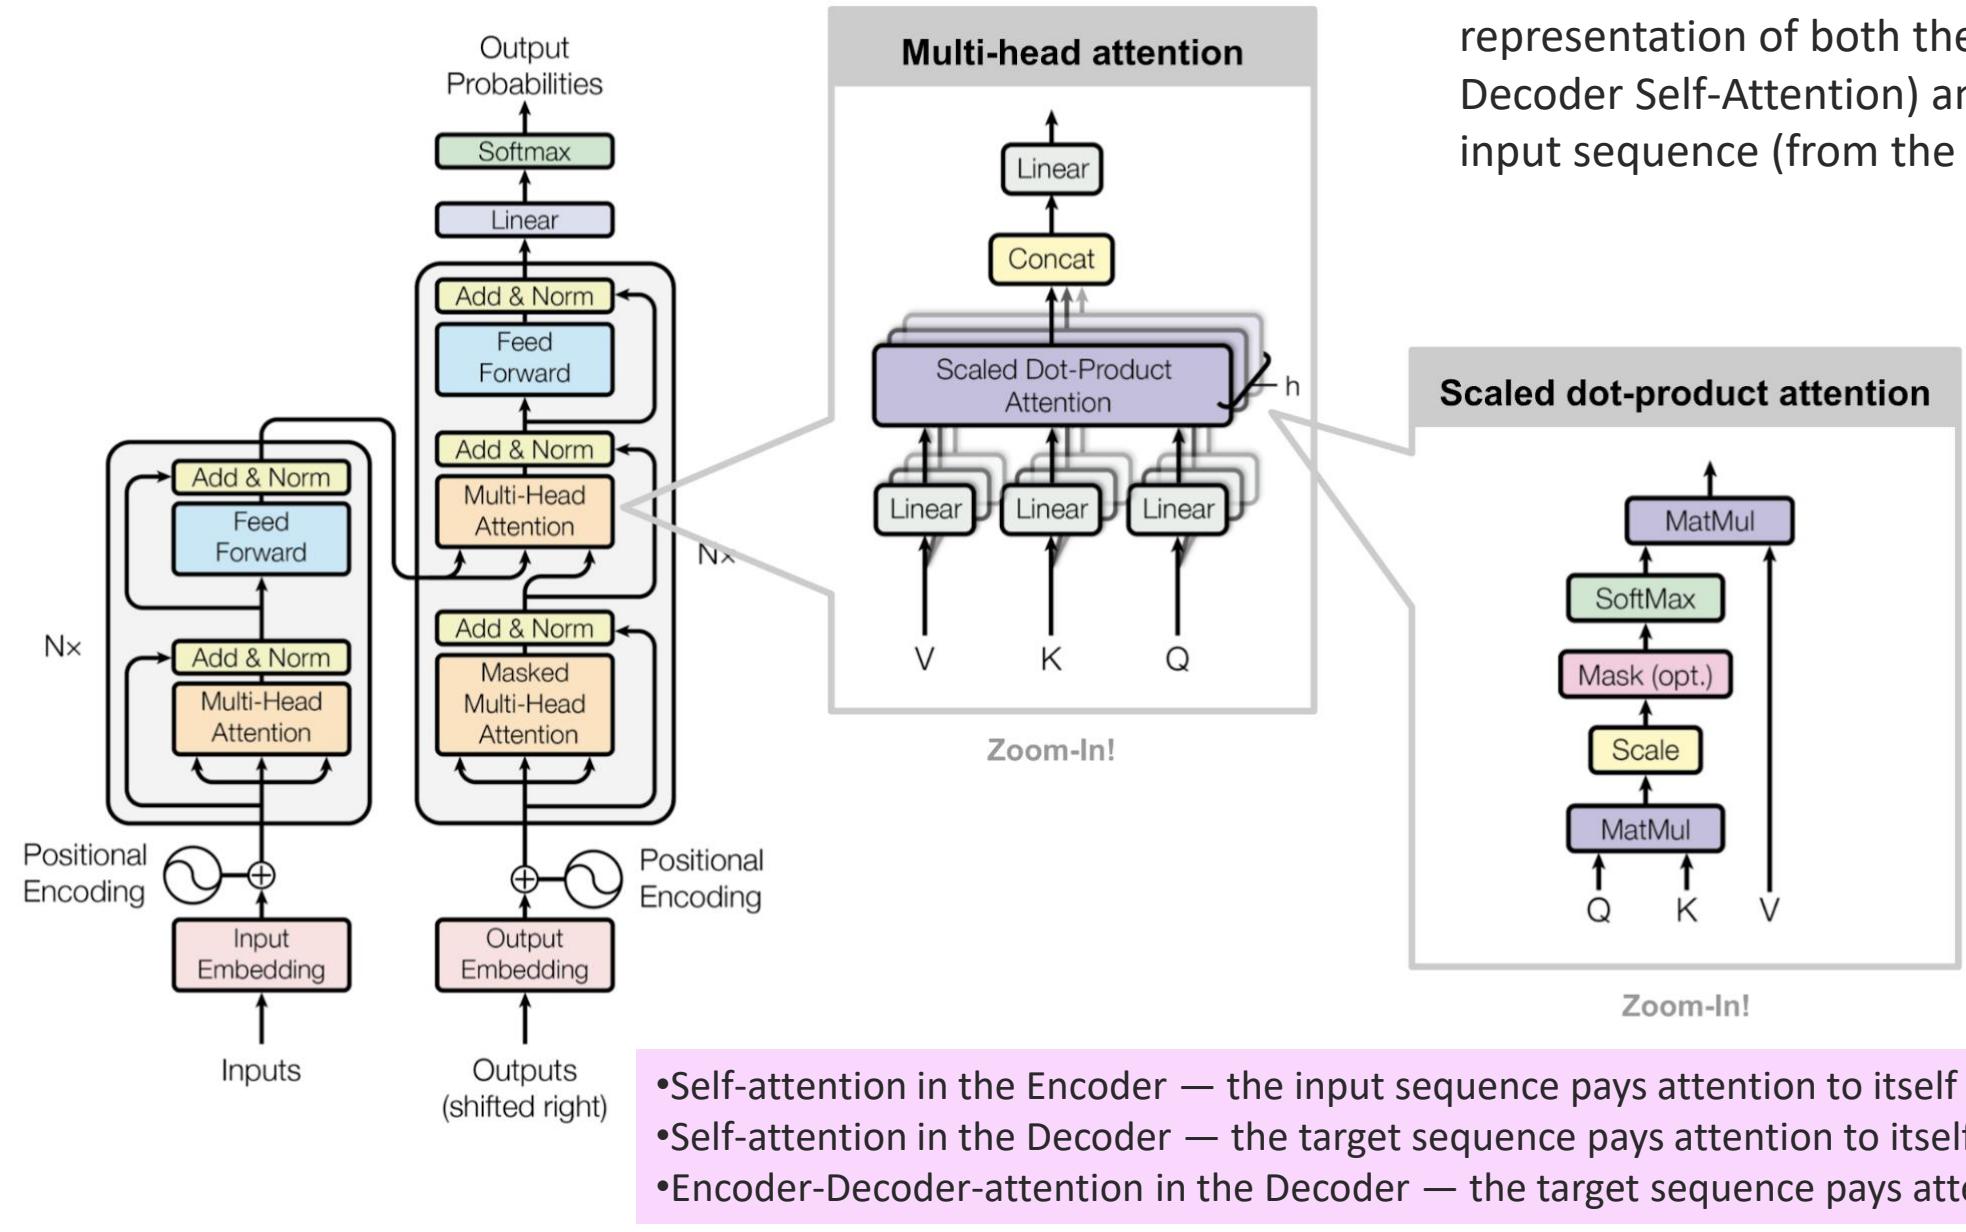
$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$
$$x = \text{LayerNorm}(x + \text{FFN}(x))$$



$$h_{ij}^{norm} = \frac{h_{ij} - \mu_j}{\sigma_j}$$

$$h_{ij}^{final} = \gamma_j \cdot h_{ij}^{norm} + \beta_j$$

# Transformers



The Encoder-Decoder Attention is getting a representation of both the target sequence (from the Decoder Self-Attention) and a representation of the input sequence (from the Encoder stack).

# Encoder –Decoder Cross Attention

- The decoder generates tokens **one by one** and, at each step, it needs to determine **which encoder outputs to attend to**.
- To do this, the decoder uses **cross-attention**, which works as follows:
  - Query ( $Q$ ) → Comes from the decoder (its current hidden state).
  - Key ( $K$ ) & Value ( $V$ ) → Come from the encoder outputs  $H_{\text{enc}}$ .
- The cross-attention mechanism computes **attention scores** using the **scaled dot-product attention**:

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

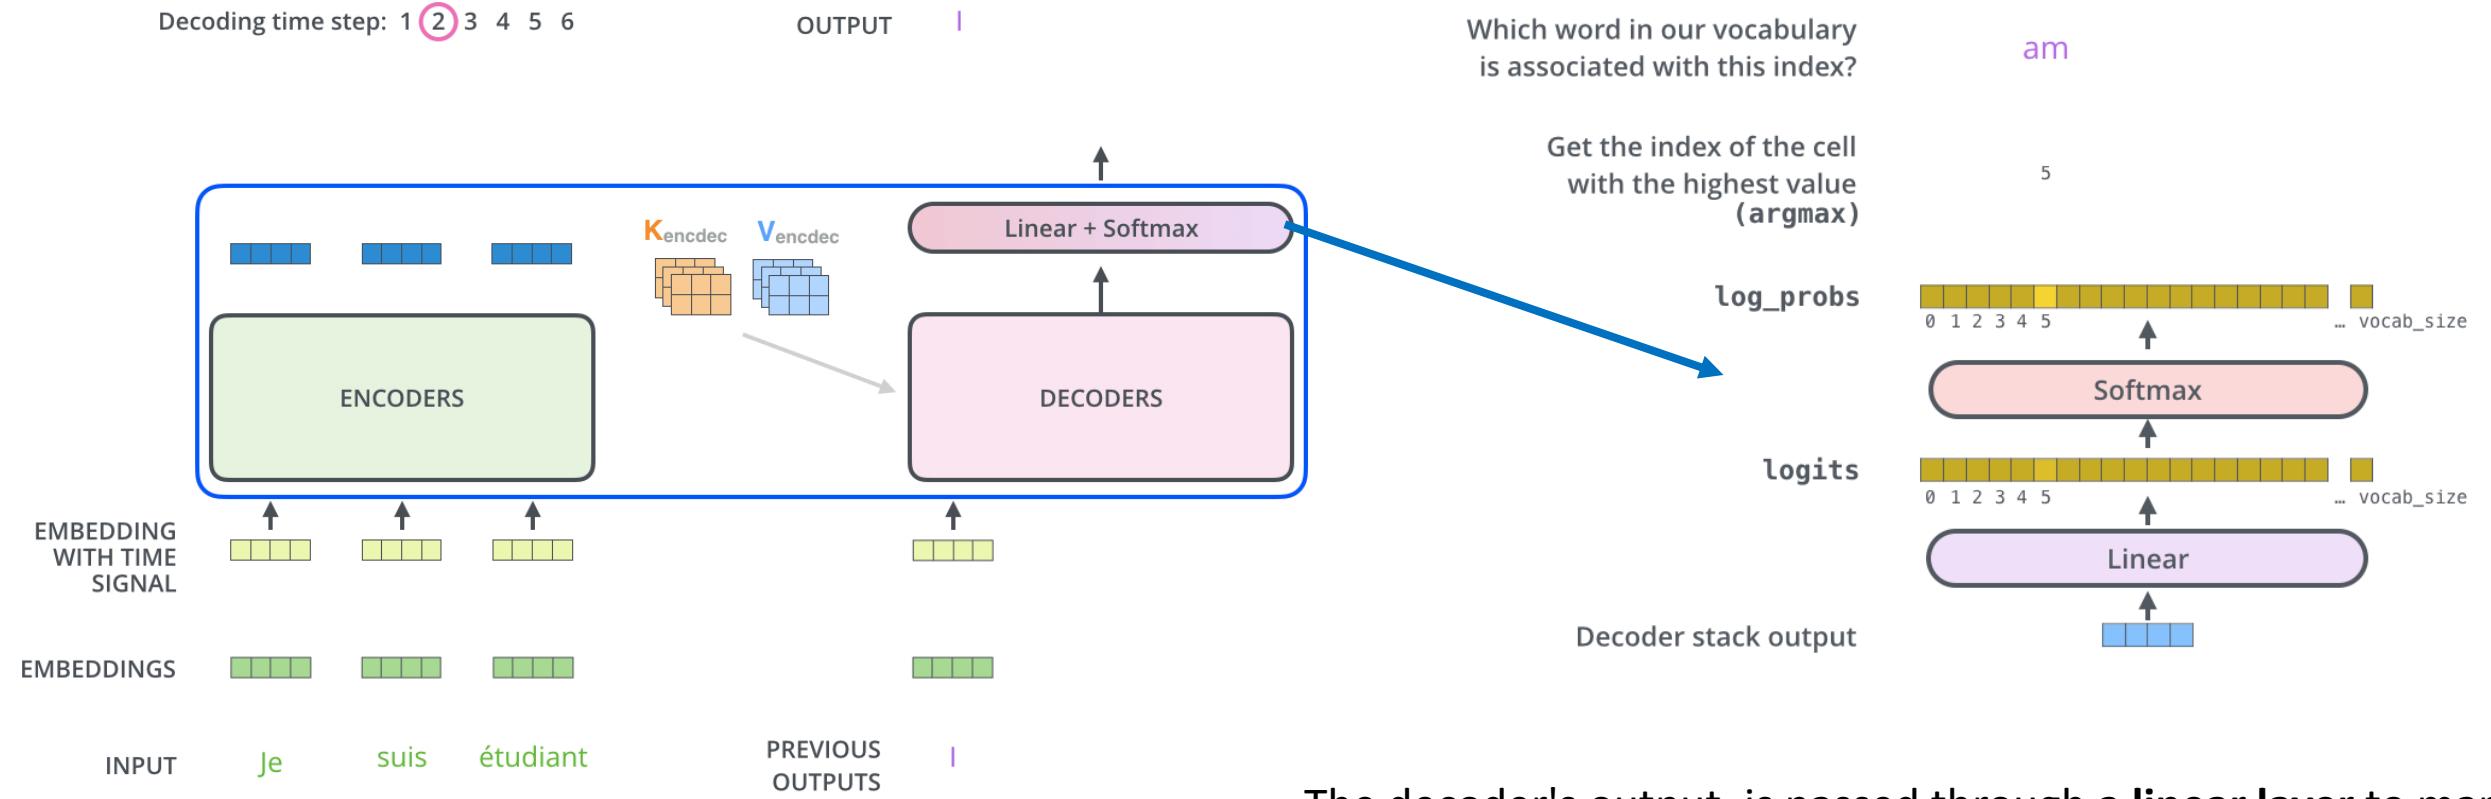
$$K = H_{\text{enc}}W_K$$

$$V = H_{\text{enc}}W_V$$

- The **dot product**  $QK^T$  measures how relevant each encoder output  $h_i$  is to the current decoder query.
- The **softmax** function normalizes these scores into attention weights.
- The final output is a **weighted sum** of the encoder representations.

# Transformer

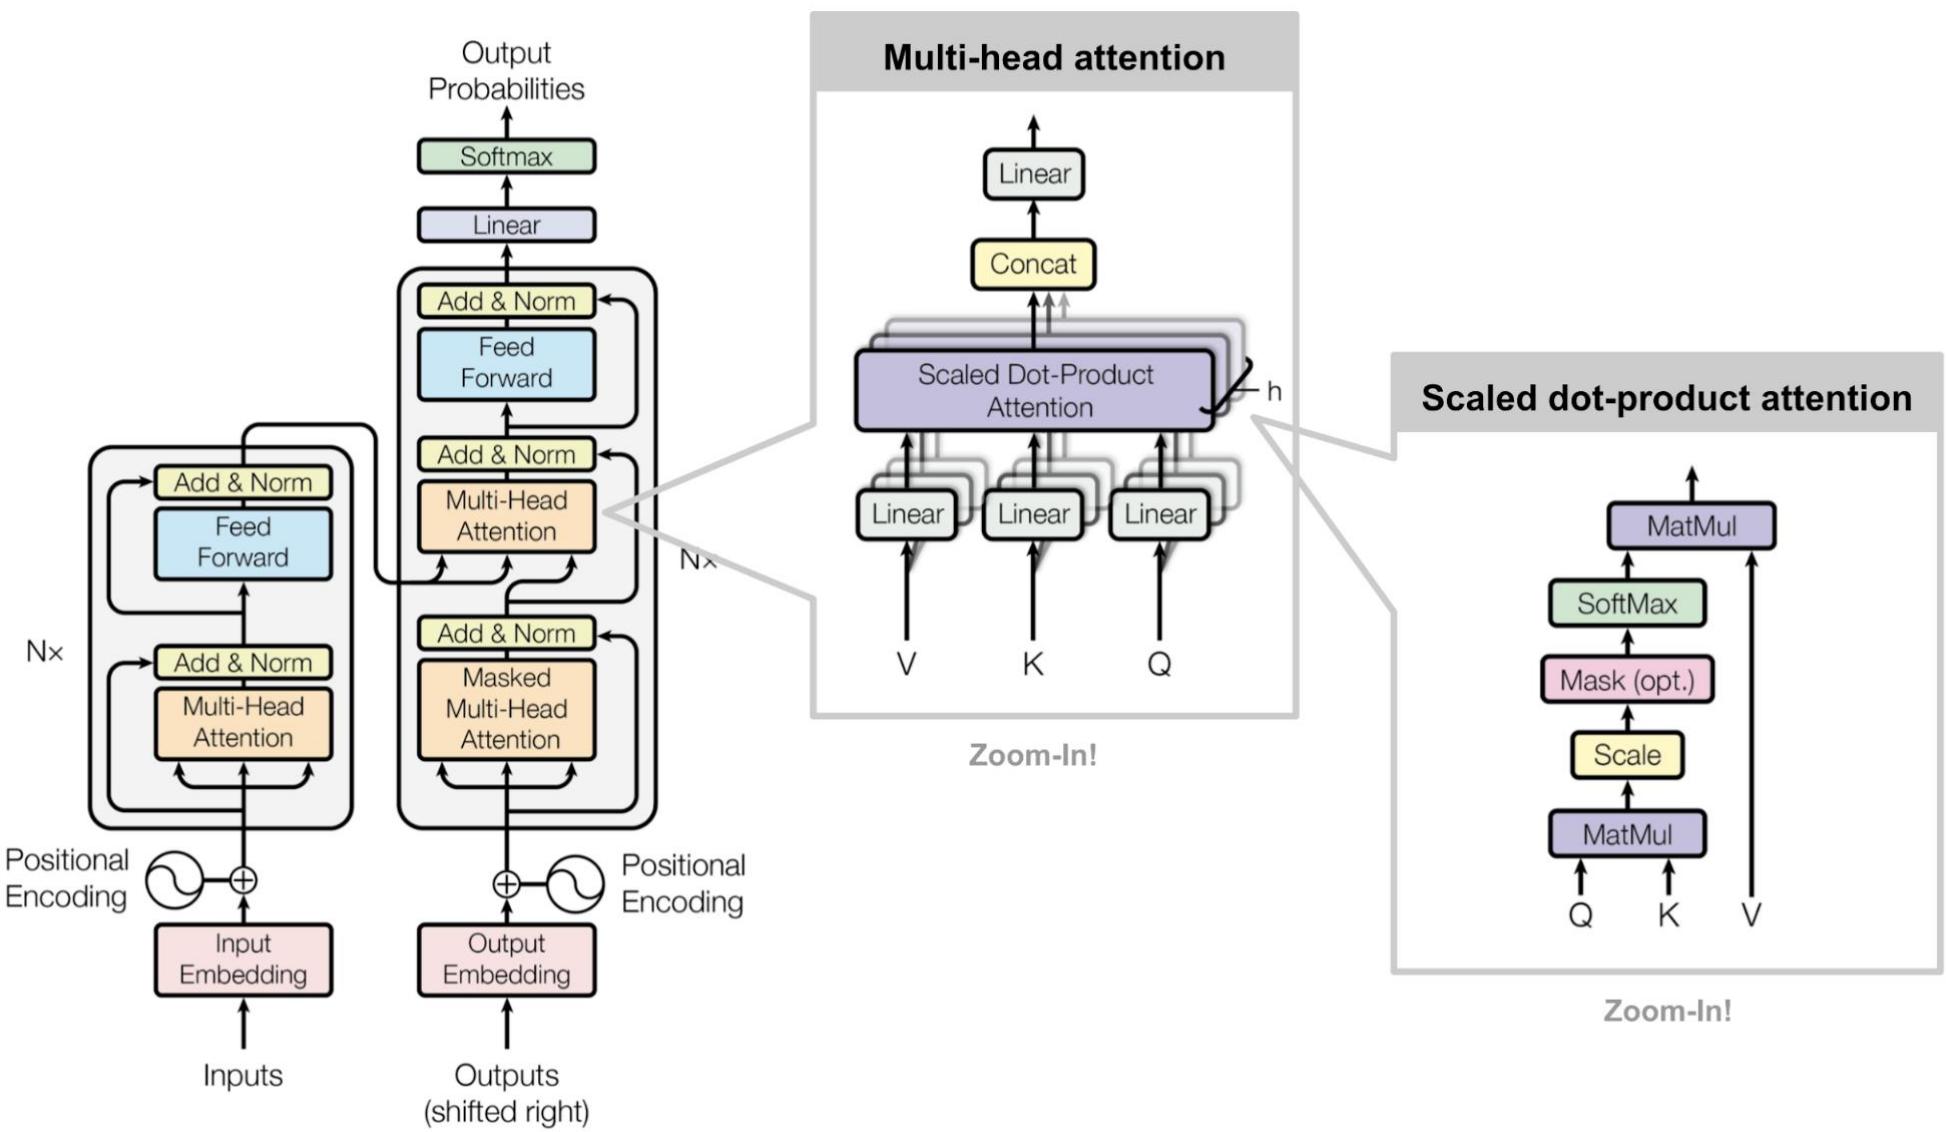
The “Encoder-Decoder Attention” layer works just like multiheaded self-attention, except it creates its Queries matrix from the layer below it, and takes the Keys and Values matrix from the output of the encoder stack.



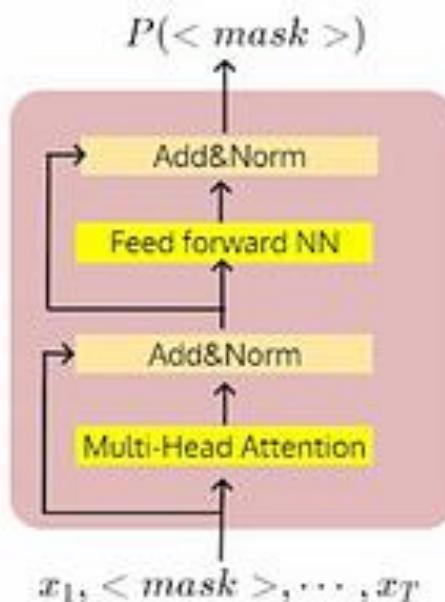
The self attention layers in the decoder operate in a slightly different way than the one in the encoder:

The decoder's output is passed through a **linear layer** to map it from **dmodel** dimensions to **vocabulary size**.

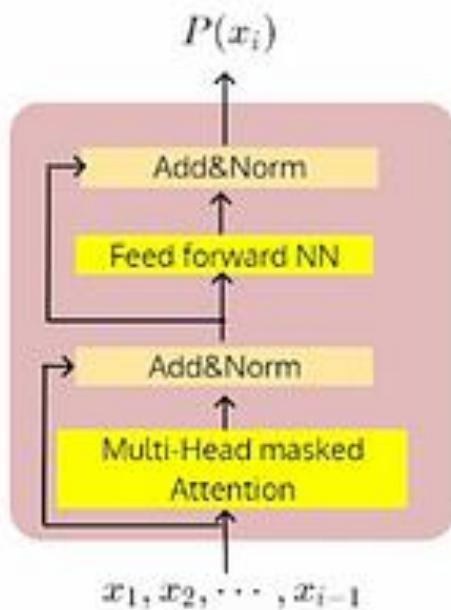
# Are you clear with all the blocks?



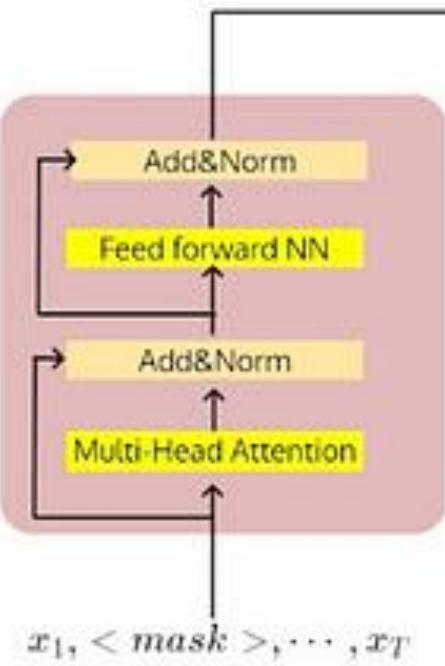
## Some Possibilities



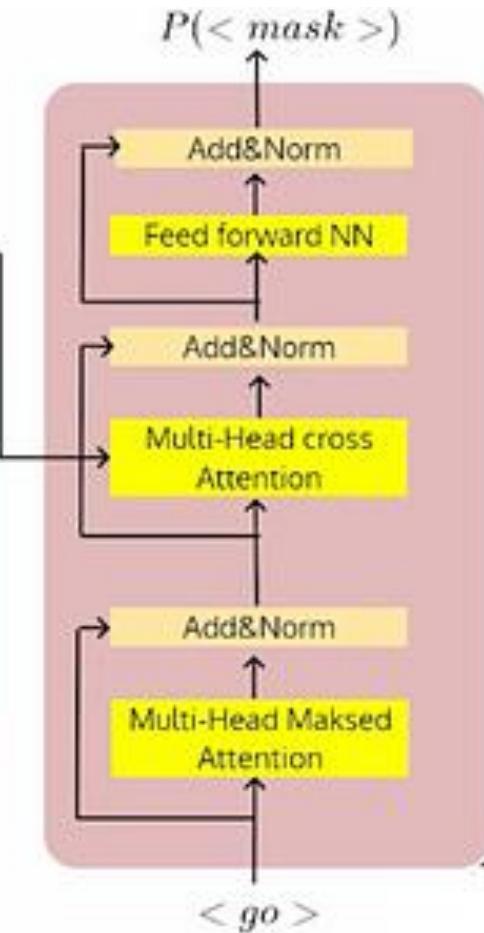
Using only the encoder of  
the transformer (encoder  
only models)

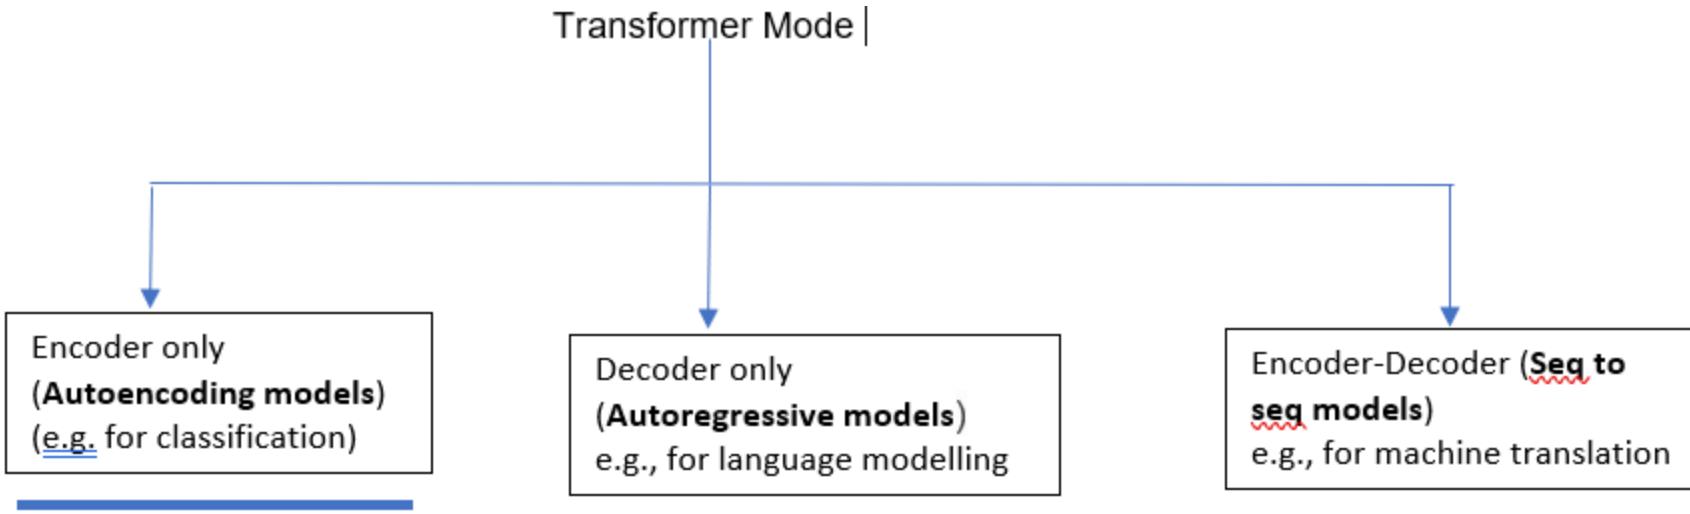


Using only the decoder of  
the transformer (decoder  
only models)



Using both the encoder and  
decoder of the transformer  
(encoder-decoder models)





### Autoencoding models

- BERT
- ALBERT
- RoBERTa
- DistilBERT
- XLM
- XLM RoBERTa
- FlauBERT
- ELECTRA
- Longformer

### Autoregressive models

- XLNET
- Reformer
- Transformer XL
- GPT
- GPT-2

### Seq-seq models

- BART
- MarianMT
- T5
- CTRL

Longformer, Linformer, BigBird, Reformer, Long Range Arena, Poolingformer, conformer, GPT 3.5, GPT 4,4.5

[View PDF](#)[Download full issue](#)[Outline](#)[Abstract](#)[Keywords](#)[1. Introduction](#)[2. Background](#)[3. Taxonomy of Transformers](#)[4. Attention](#)[5. Other module-level modifications](#)[6. Architecture-level variants](#)[7. Pre-trained Transformers](#)[8. Applications of Transformer](#)[9. Conclusion and future directions](#)[References](#)CHINESE ROOTS  
GLOBAL IMPACT

AI Open

Volume 3, 2022, Pages 111-132



# A survey of transformers

[Tianyang Lin, Yuxin Wang, Xiangyang Liu, Xipeng Qiu](#) [Show more ▾](#)[+ Add to Mendeley](#) [Share](#) [Cite](#)<https://doi.org/10.1016/j.aiopen.2022.10.001> ↗[Get rights and content ↗](#)[Under a Creative Commons license ↗](#) [open access](#)

## Abstract

Transformers have achieved great success in many artificial intelligence fields, such as [natural language processing](#), computer vision, and [audio processing](#). Therefore, it is natural to attract lots of interest from academic and industry researchers. Up to the

[Recommended art](#)[Causes of transformer diagnostic methods](#)

Renewable and Sustainable  
Christina Aj, ..., William V

[Purchase PDF](#)[Self-directed machir](#)

AI Open, Volume 3, 2022,  
Wenwu Zhu, ..., Pengtao J

[View PDF](#)[Hierarchical label with attributed networks](#)

AI Open, Volume 3, 2022,  
Shu Zhao, ..., Jie Tang

[View PDF](#)[Show 3 more articles ↗](#)

## Computer Science &gt; Computation and Language

[Submitted on 12 Feb 2023 (v1), last revised 16 Feb 2023 (this version, v2)]

# Transformer models: an introduction and catalog

Xavier Amatriain

In the past few years we have seen the meteoric appearance of dozens of models of the Transformer family, all of which have funny, but not self-explanatory, names. The goal of this paper is to offer a somewhat comprehensive but simple catalog and classification of the most popular Transformer models. The paper also includes an introduction to the most important aspects and innovation in Transformer models.

Subjects: **Computation and Language (cs.CL)**

Cite as: arXiv:2302.07730 [cs.CL]

(or arXiv:2302.07730v2 [cs.CL] for this version)

<https://doi.org/10.48550/arXiv.2302.07730> 

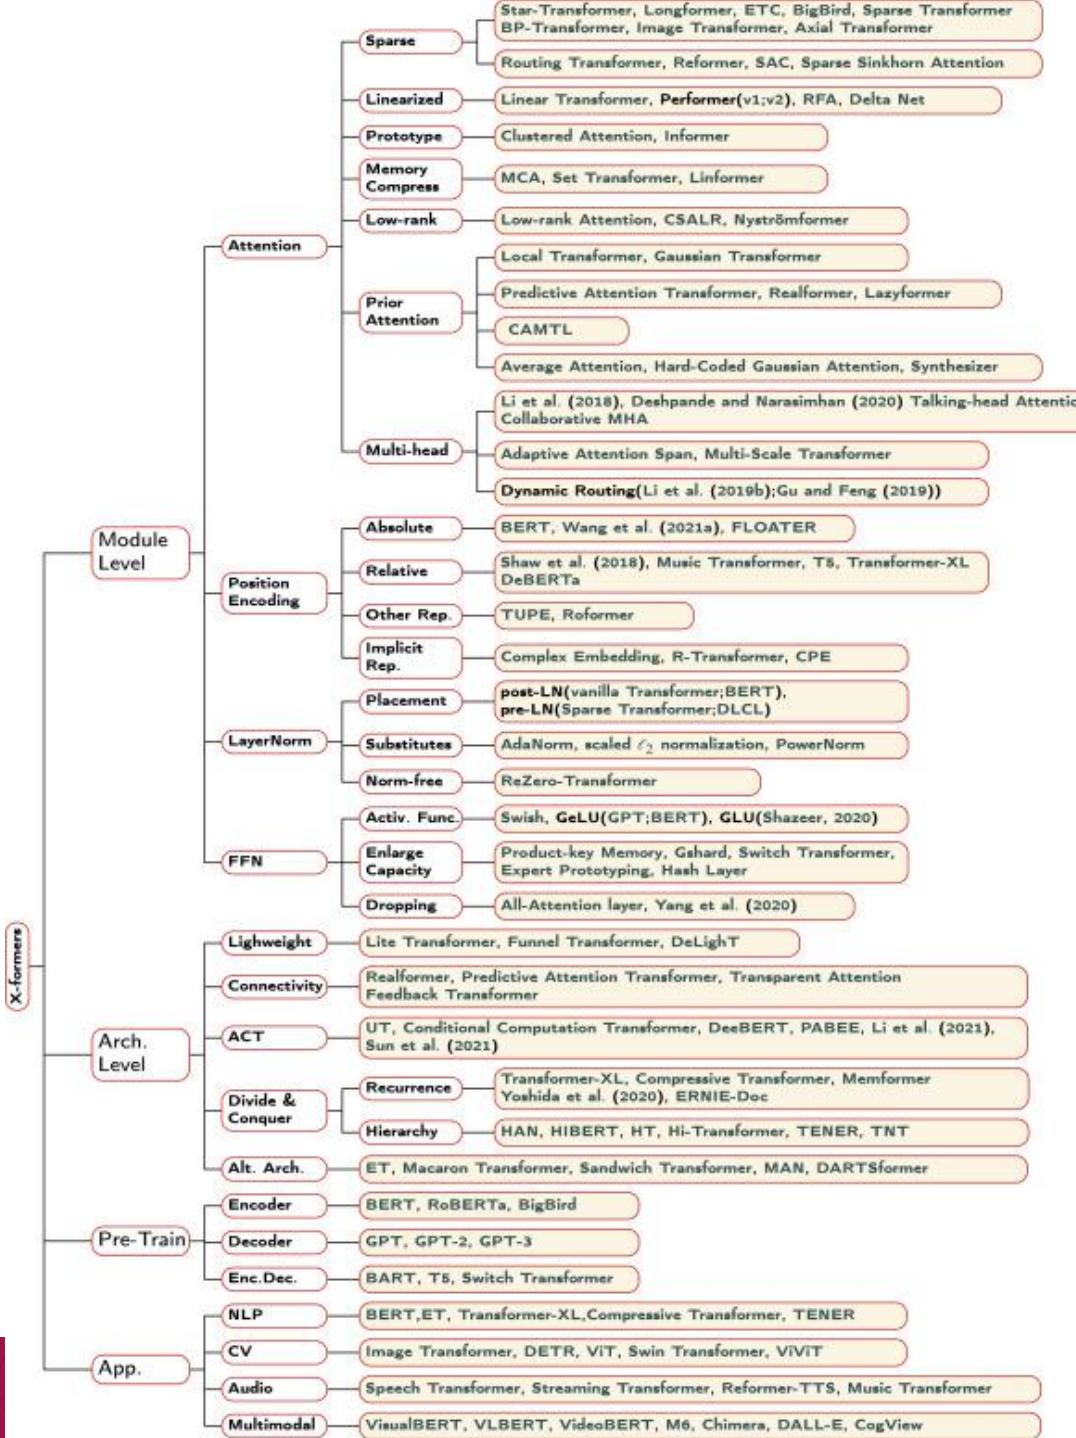
## Submission history

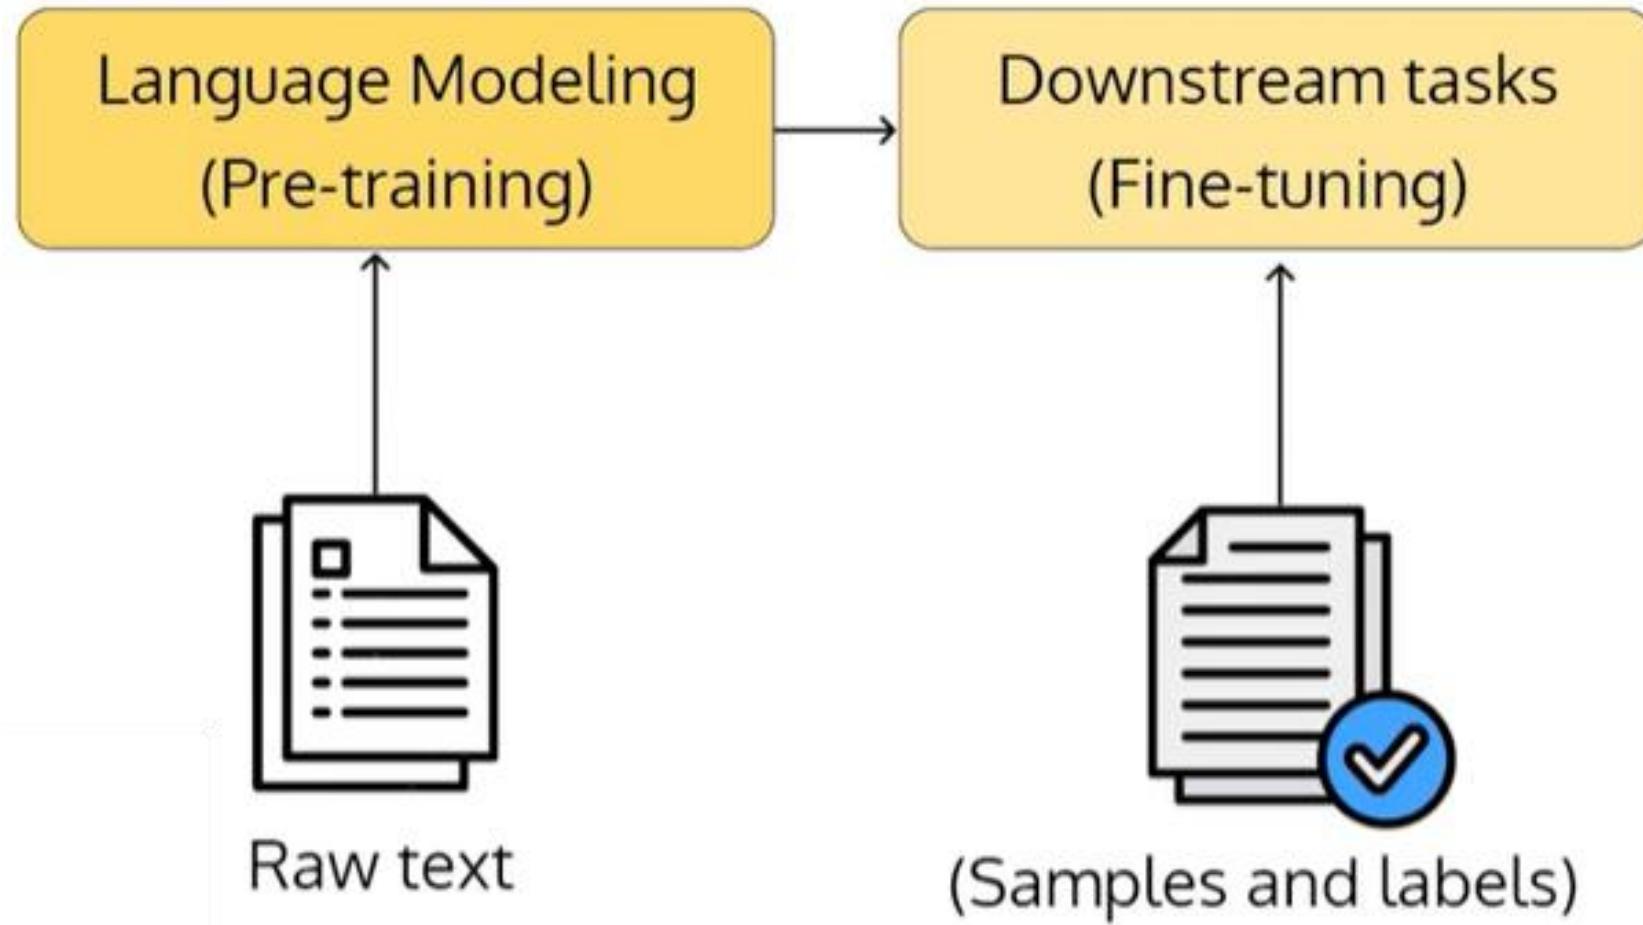
From: Xavier Amatriain [[view email](#)]

[v1] Sun, 12 Feb 2023 01:26:49 UTC (3,568 KB)

[v2] Thu, 16 Feb 2023 05:31:15 UTC (3,568 KB)

<https://arxiv.org/pdf/2302.07730v2.pdf>





# Down-stream task

What if we want to use the transformer architecture for other NLP tasks?

We need to train a separate model for each task using dataset specific to the task

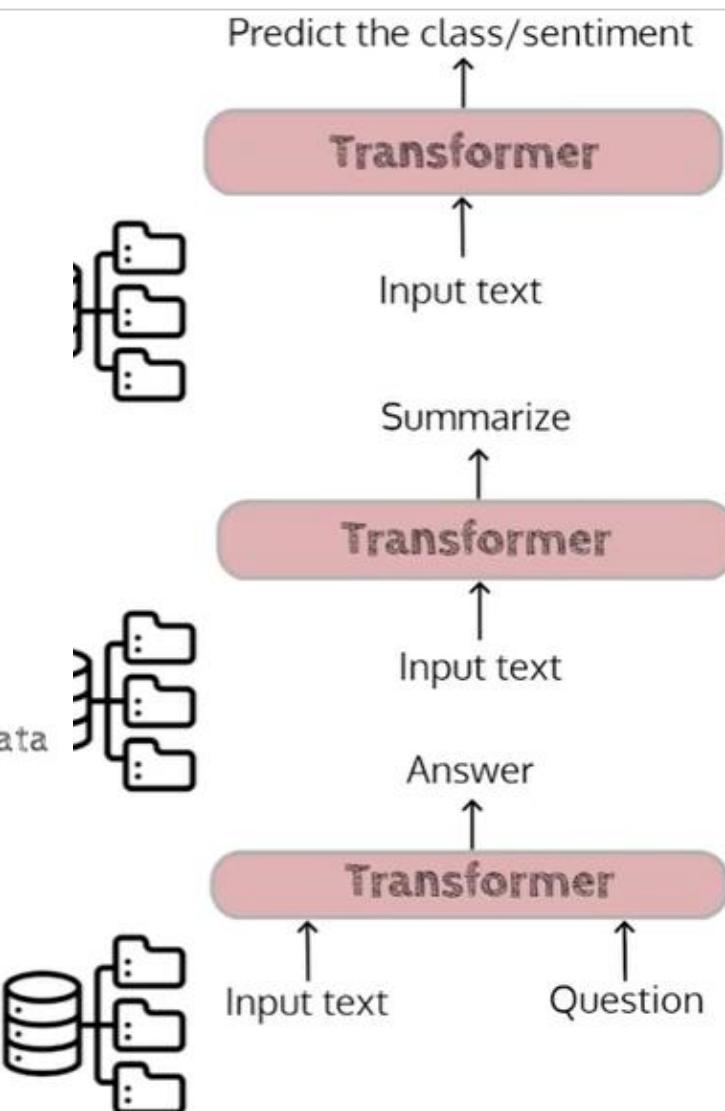
If we train the architecture from scratch (that is, by randomly initializing the parameters) for each task, it takes a long time for convergence

Often, we may not have enough **labelled** samples for many NLP tasks

We have a large amount of unlabelled text easily available on the internet



Can we make use of such unlabelled data to train a model?



Assume that we ask **questions** to a lay person based on a statement or some excerpt

"Wow, India has now reached the moon"

An excerpt from business today "What sets this mission apart is the pivotal role of artificial intelligence (AI) in guiding the spacecraft during its critical descent to the moon's surface."

Is this sentence expressing a positive or a negative sentiment?

Did the lander use AI for soft landing on the moon?

He likes to stay  
He likes to stray  
He likes to sway

Are these meaningful sentences?

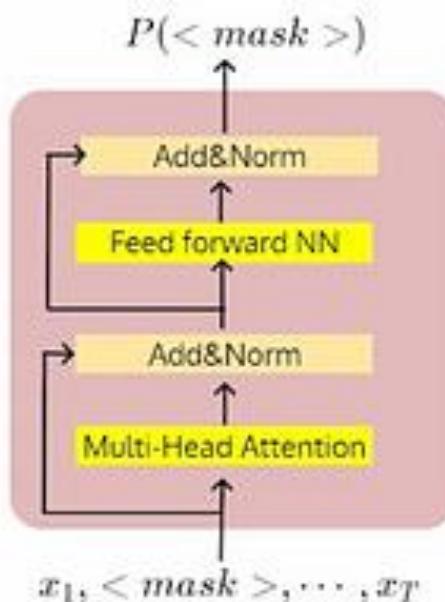
The person will most likely answer all the questions, even though he/she may not be explicitly trained on any of these tasks. How?

We develop a strong understanding of language through various language based interactions( listening/reading) over our life time without any explicit supervision

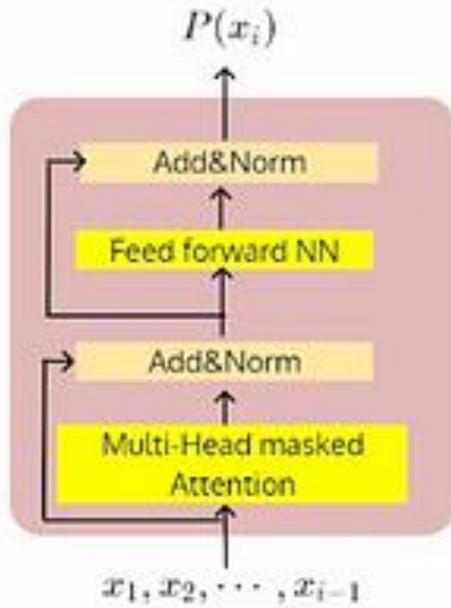
Can a model develop basic understanding of language by getting exposure to a large amount of raw text? [pre-training]

More importantly, after getting exposed to such raw data can it learn to perform well on downstream tasks with minimum supervision? Supervised fine tuning

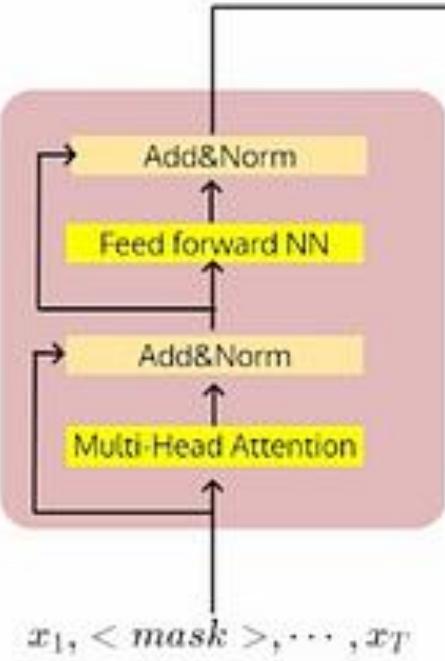
## Some Possibilities



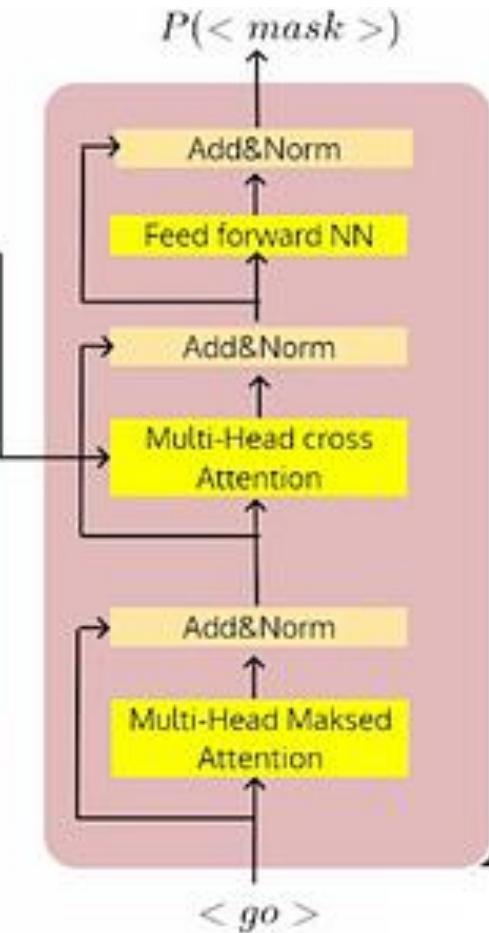
Using only the encoder of the transformer (encoder only models)



Using only the decoder of the transformer (decoder only models)



Using both the encoder and decoder of the transformer (encoder-decoder models)



# Namah Shivaya